

Changing random stuff until your program works is
“hacky” and “bad coding practice”.
But if you do it fast enough, it’s “Machine Learning” and
pays 4x your current salary

- Unknown

Convolutional Neural Networks

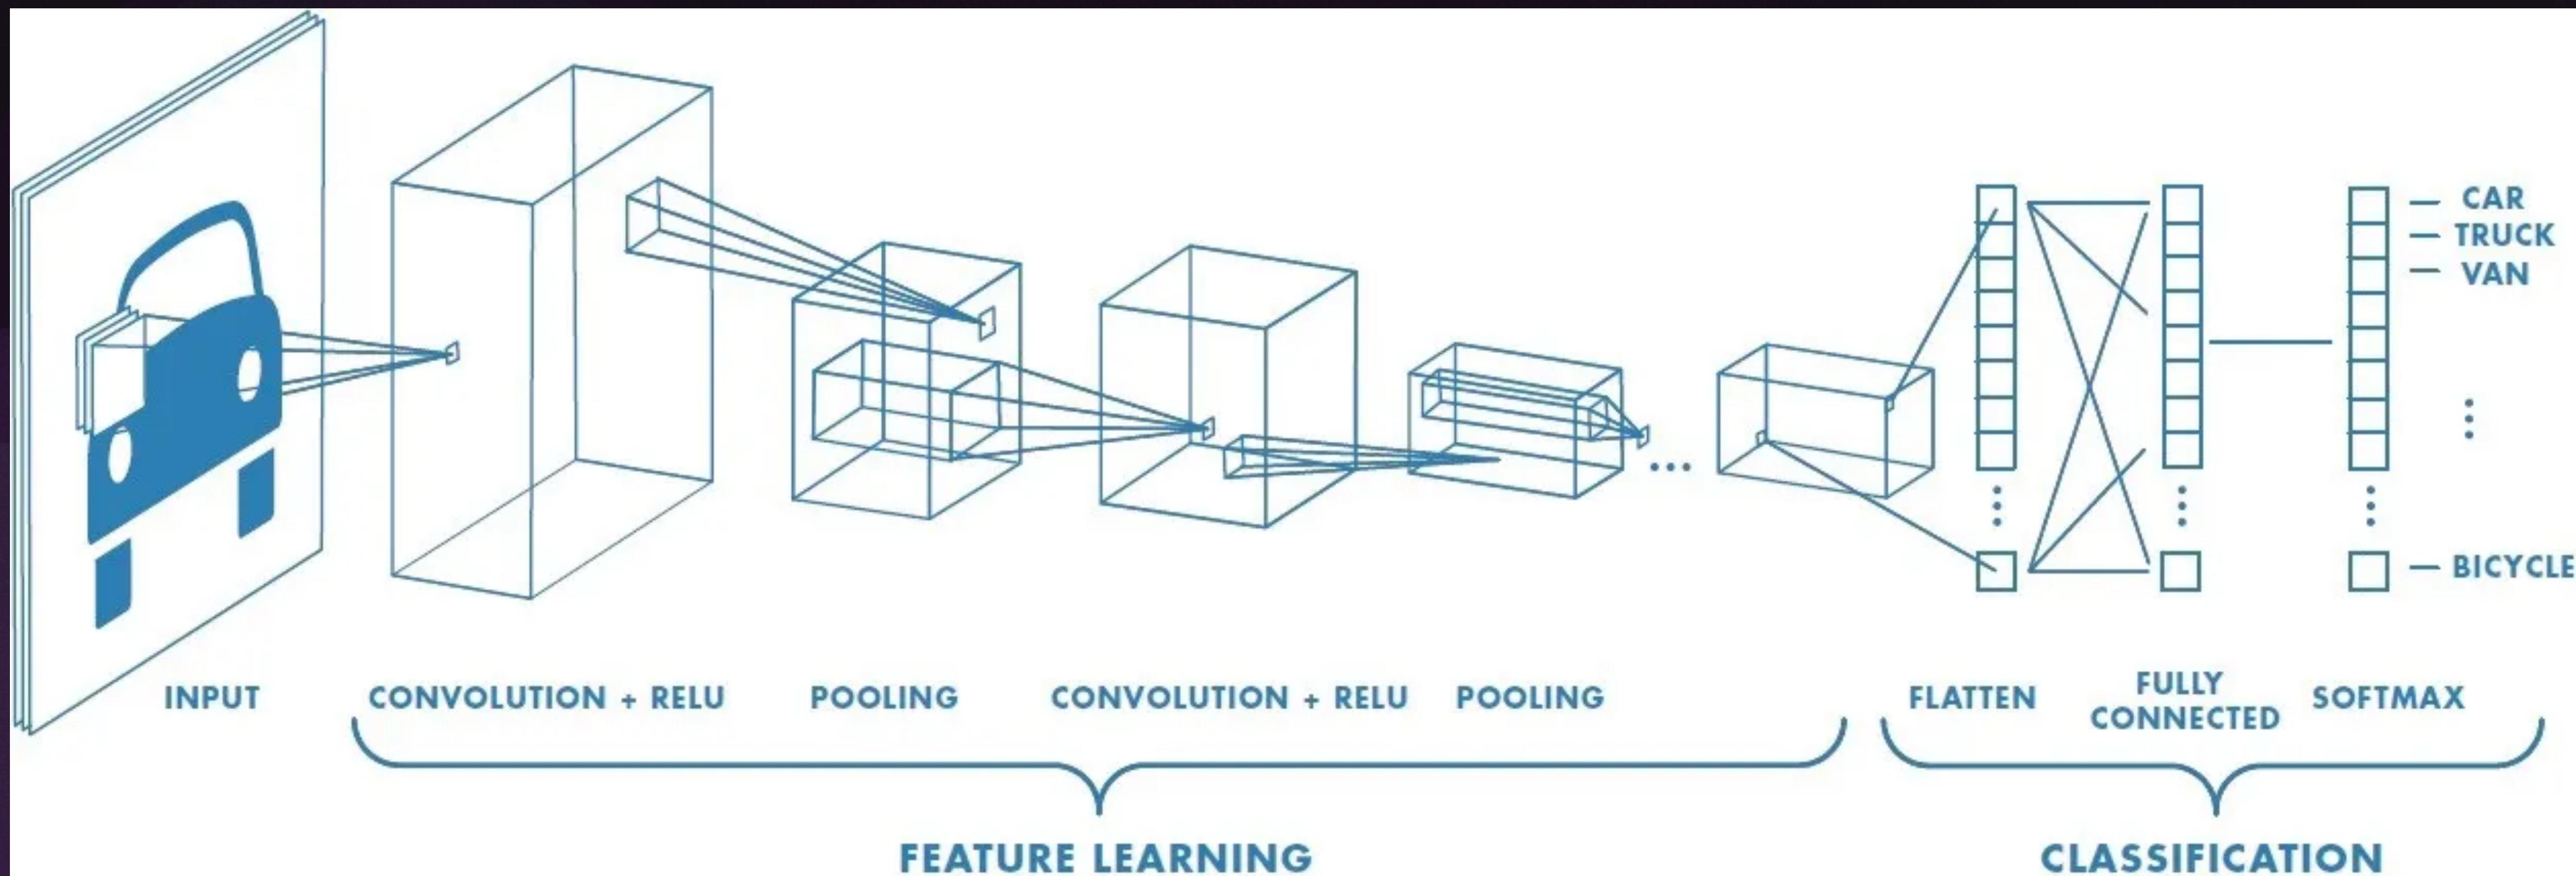
AI Systems Implementation

George Langrudi
Based on slides by Gareth Howells

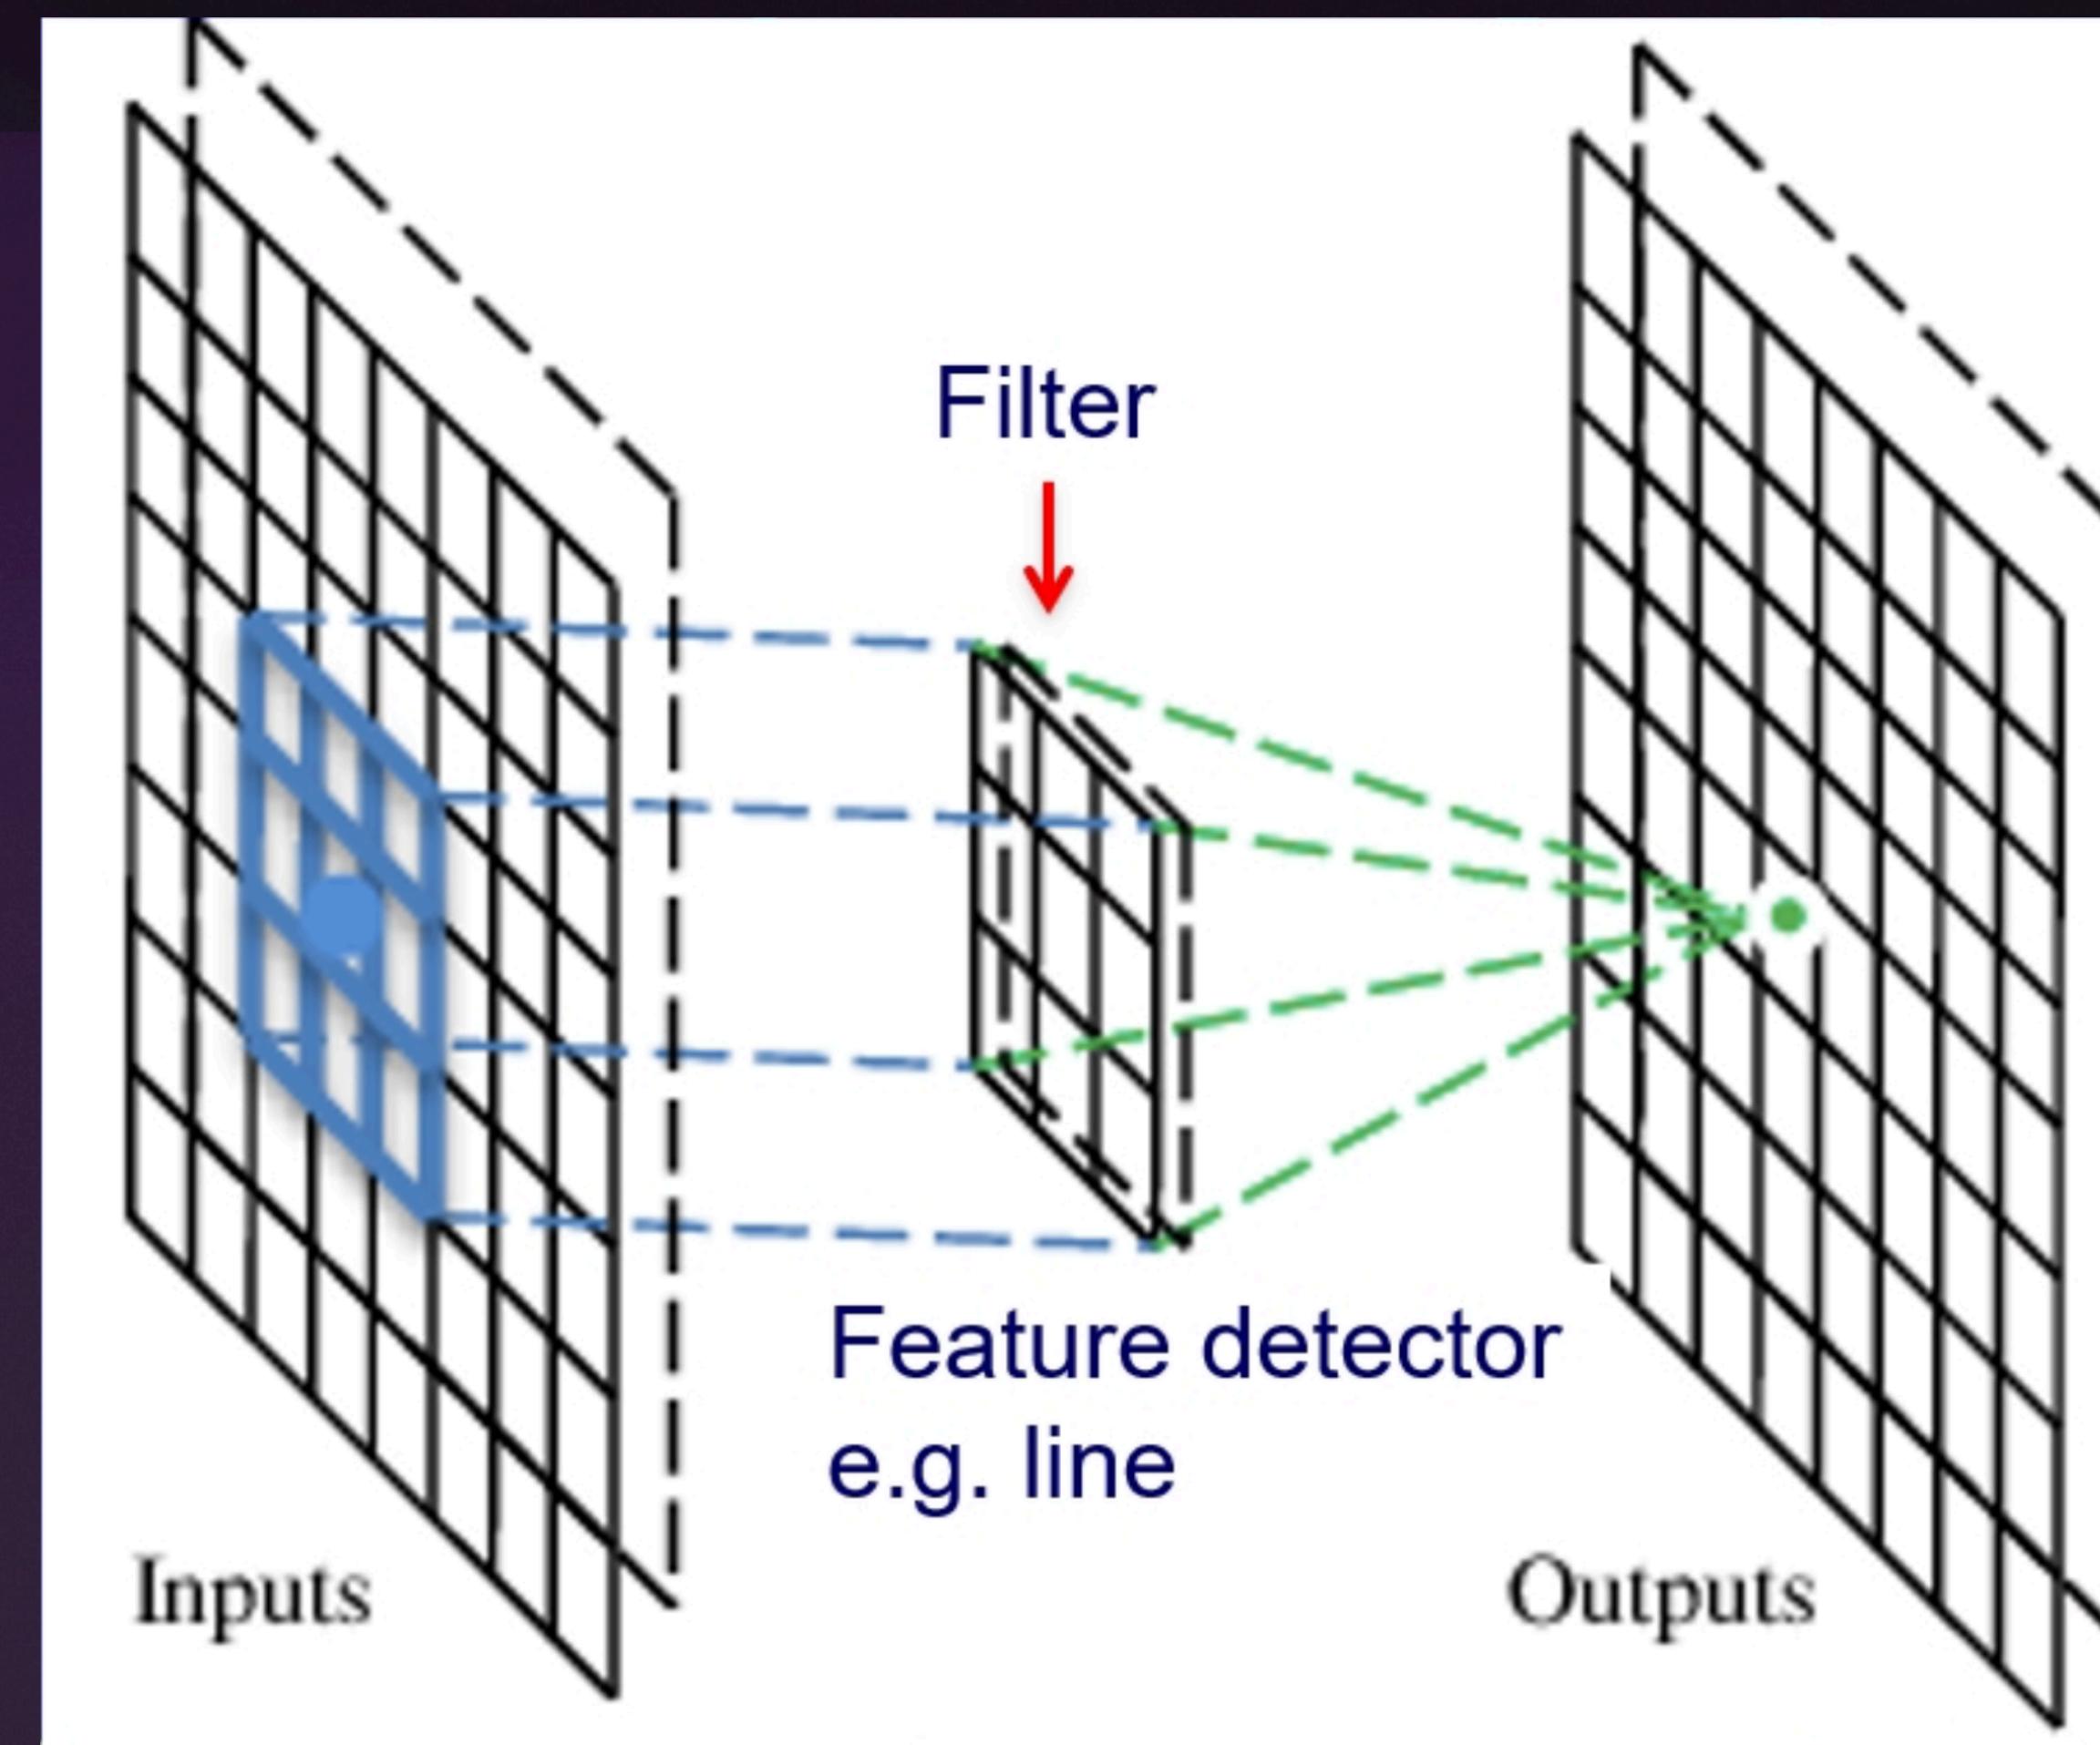
Overview

- Convolutional Neural Networks
 - Architecture
 - Filters / Convolution
 - Pooling
 - Flattening
- Fully Connected Layers

Convolutional Neural Networks

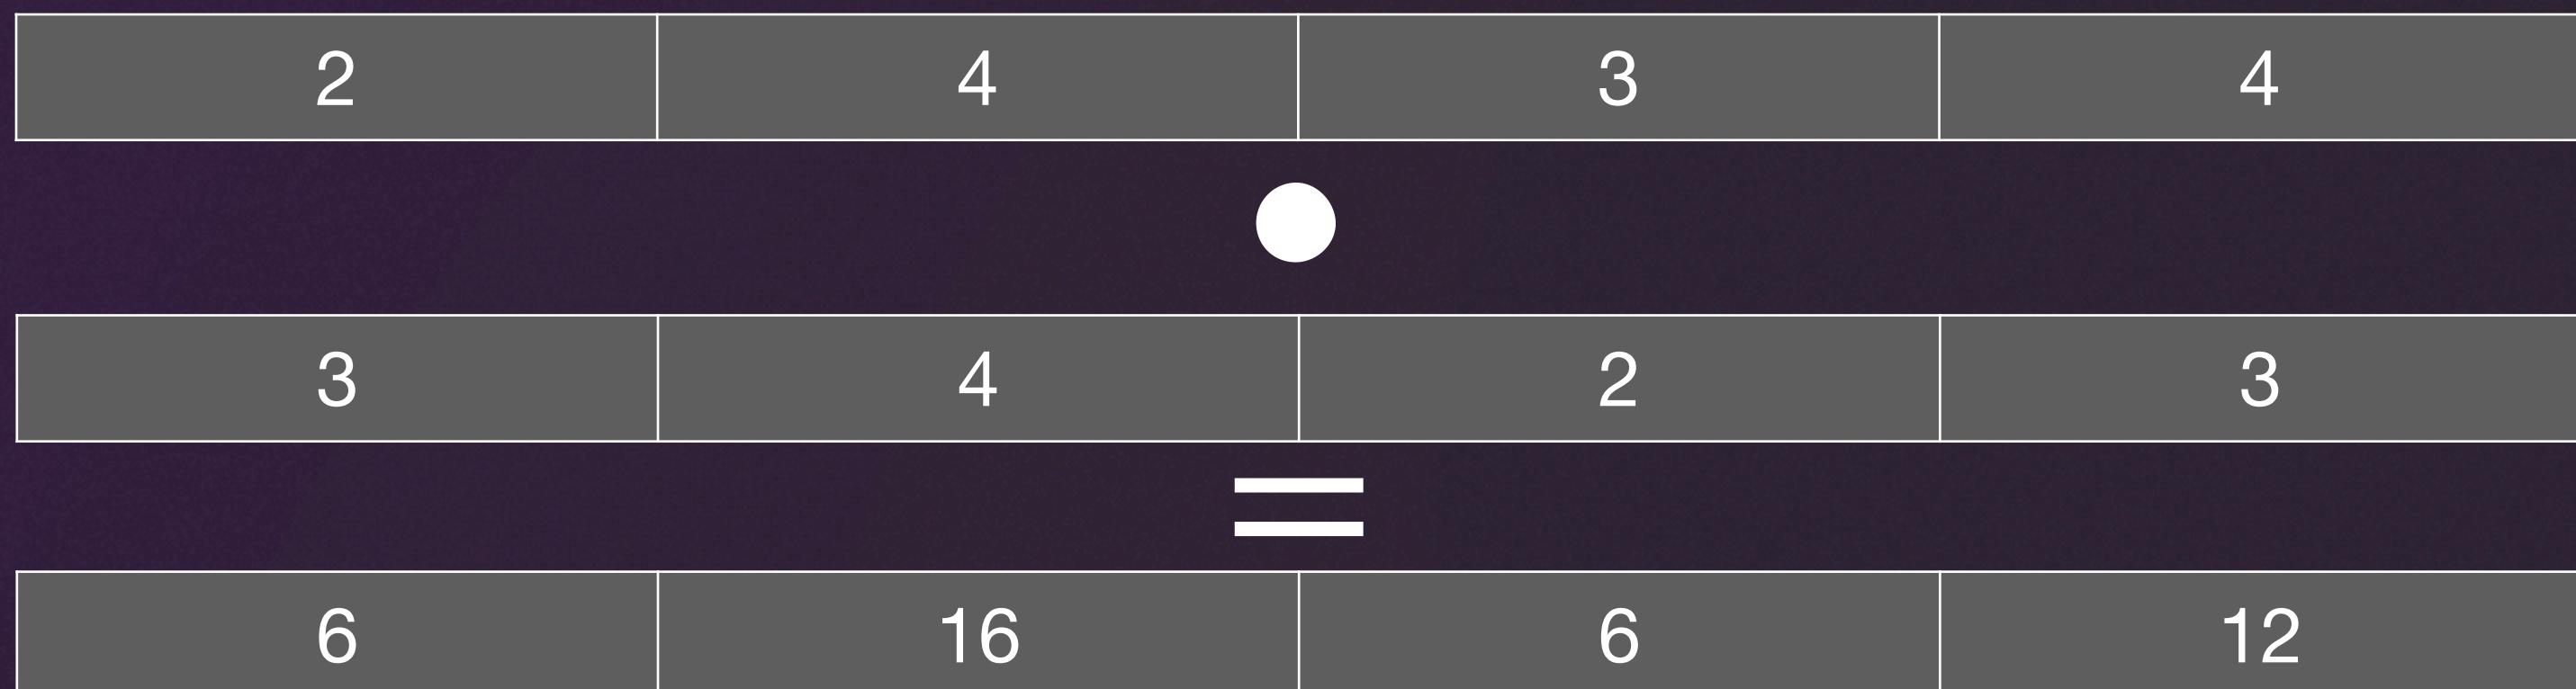


A Convolutional layer



Wait, what is convolution?

- First, we need to understand *dot product*
- $$\text{dotproduct}_{ab} = \sum_{i=1}^n a_i b_i$$
- *Essentially*, multiple every point in a by it's corresponding point in b



Wait, what is convolution?

- $(a * b)_k = \sum_{i=1}^n a_i b_{k-i}$
- $(a * b)_k$ = The convolution of a and b , at point k
- $\sum_{i=1}^n a_i b_{k-i}$ = Calculating the dot product at point k
- OK... what?

1-Dimensional Convolution

$$(1 \times 4) + ((2 \times 3)) + ((3 \times 2)) + (23 \times 4) + (3 \times 1)$$

1	2	3
---	---	---

4	3	2	4	1	5	3	5
---	---	---	---	---	---	---	---

16	19	13	21	20	26
----	----	----	----	----	----

2 Dimensional Convolution

Filters

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

-1	0	1
-2	0	2
-1	0	1

Vertical lines

-1	-2	-1
0	0	0
1	2	1

Horizontal lines

2 Dimensional Convolution

Filters

Stride = 1

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

6×6 image

4	2	-1	0
-4	-3	-1	1
-4	-3	0	2
0	0	0	0

-1	-2	-1
0	0	0
1	2	1

Sobel Horizontal filter

2 Dimensional Convolution

Filters

Stride = 1

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

6×6 image

0	-2	-1	0
0	-3	-1	3
0	-1	2	4
0	0	4	4

-1	0	1
-2	0	2
-1	0	1

Vert Horizontal filter

Example



The shape of the filter allows us to extract spatial features from our data

Lets use Sobel filters to do edge detection

-1	0	1
-2	0	2
-1	0	1

Vertical Sobel filter

-1	-2	-1
0	0	0
1	2	1

Horizontal Sobel filter

Example



Original Image: Black and White

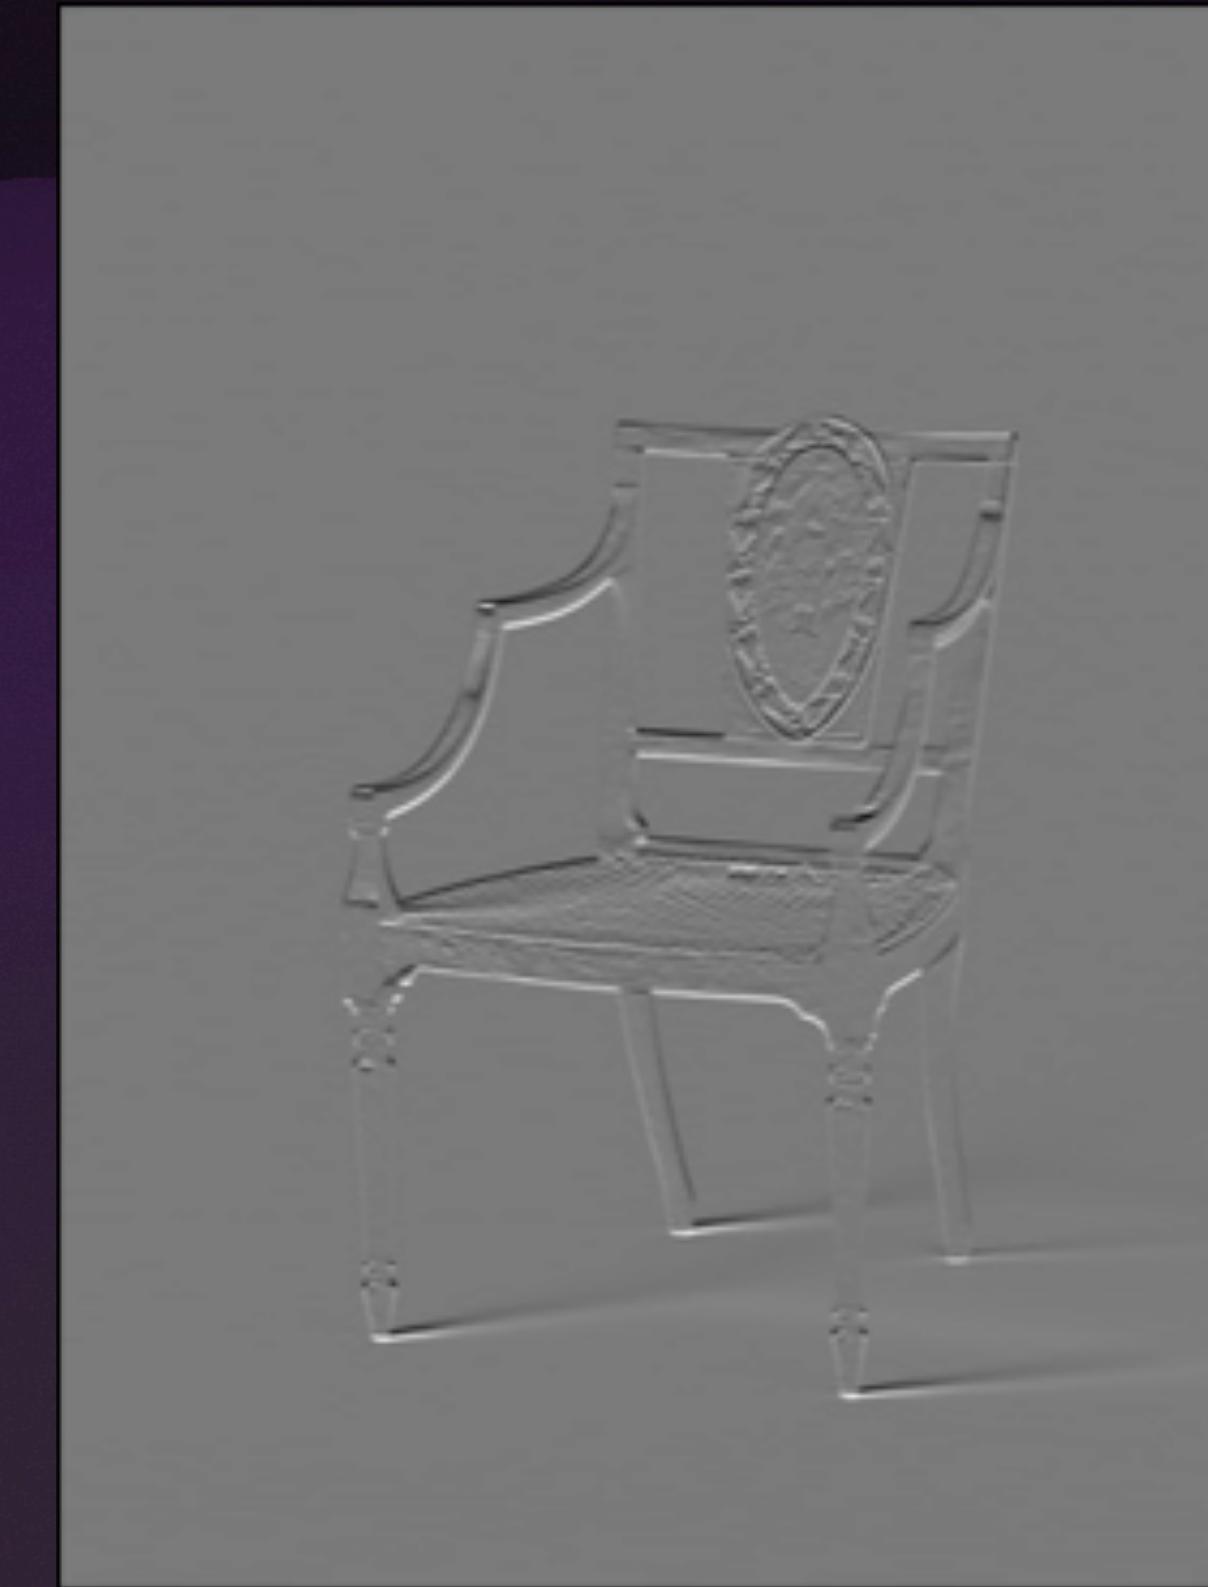


Image after convolution with a horizontal
Sobel Filter

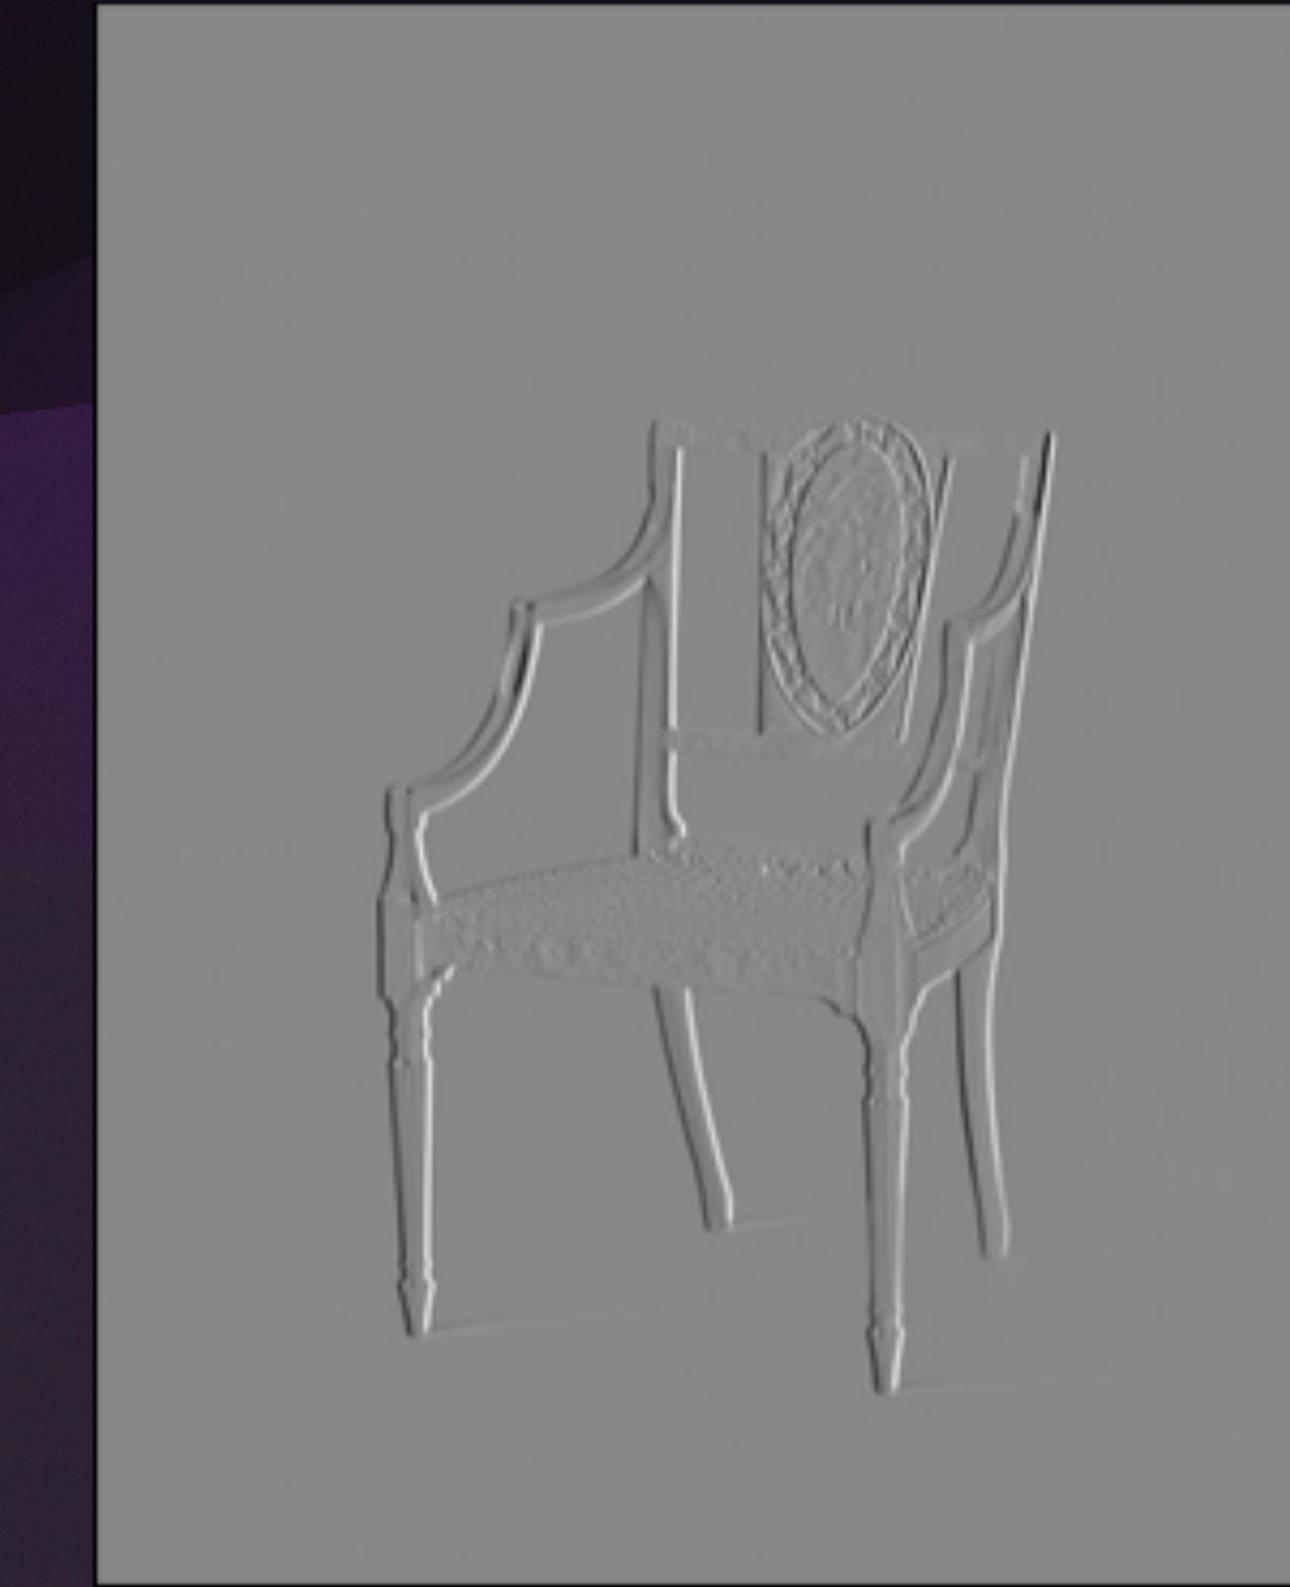


Image after convolution with a vertical
Sobel Filter

Example



```
1 def conv2d(image: np.array, kernel: np.array) -> np.array:
2     f_width, f_height = kernel.shape
3     conv_array = np.zeros([image.shape[0]+1 - f_width, image.shape[1]+1 - f_height])
4     c_width, c_height = conv_array.shape
5
6     for x in range(c_width):
7         for y in range(c_height):
8             target = image[x:x+f_width, y:y+f_height]
9             conv_array[x,y] = np.sum(conv(target, kernel))
10    return conv_array
11
12 def conv(target, filter):
13     output = np.zeros(target.shape)
14     for x in range(output.shape[1]):
15         for y in range(output.shape[1]):
16             output[x,y] = target[x,y] * filter[x,y]
17     return output
```

Pooling Layer

Max Pooling

-1	-2	-1
0	0	0
1	2	1

Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

Sobel
Vertical
filter

4	2	-1	0
-4	-3	-1	1
-4	-3	0	2
0	0	0	0

Horizontal
filtered
image

0	-2	-1	0
0	-3	-1	3
0	-1	2	4
0	0	4	4

Vertical
filtered
image

Pooling Layer

Max Pooling

-1	-2	-1
0	0	0
1	2	1

Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

Sobel
Vertical
filter

4	2
-4	-3

-1	0
-1	1

Horizontal
filtered
image

-4	-3
0	0

0	2
0	0

0	-2	-1	0
0	-3	-1	3
0	-1	2	4
0	0	4	4

Vertical
filtered
image

Pooling Layer

Max Pooling

-1	-2	-1
0	0	0
1	2	1

Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

Sobel
Vertical
filter

4	2
-4	-3

-1	0
-1	1

Horizontal
filtered
image

-4	-3
0	0

0	2
0	0

0	-2
0	-3

-1	0
-1	3

Vertical
filtered
image

0	-1
0	0

2	4
4	4

Pooling Layer

Max Pooling

-1	-2	-1
0	0	0
1	2	1

Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

Sobel
Vertical
filter

4

1

Horizontal
filtered
image

-4	-3
0	0

0	2
0	0

0	-2
0	-3

-1	0
-1	3

Vertical
filtered
image

0	-1
0	0

2	4
4	4

Pooling Layer

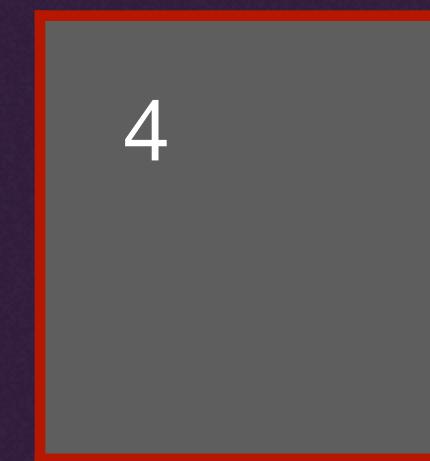
Max Pooling

-1	-2	-1
0	0	0
1	2	1

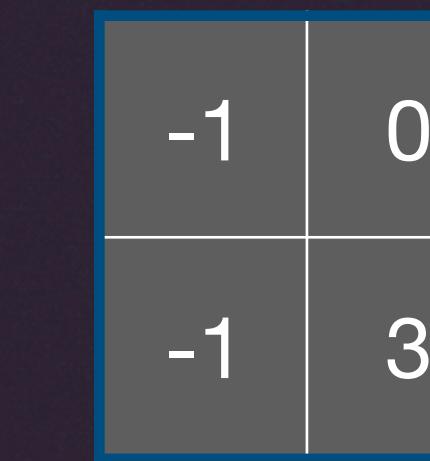
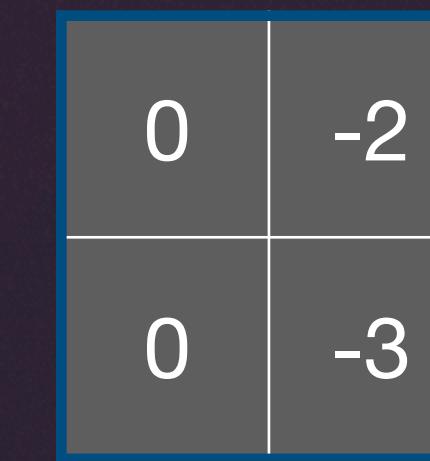
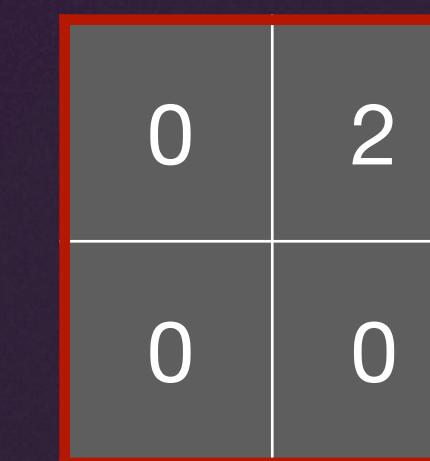
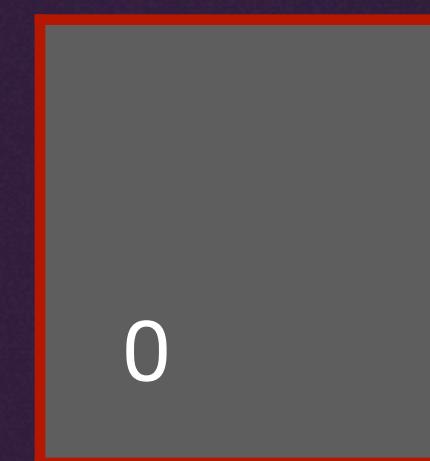
Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

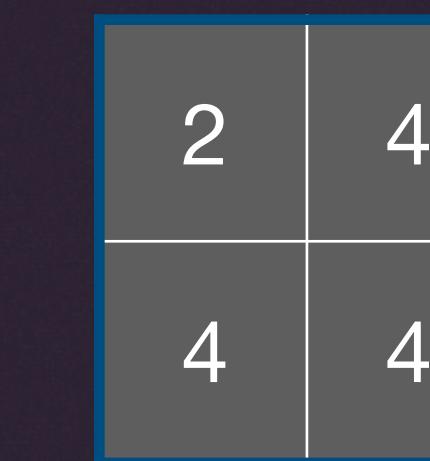
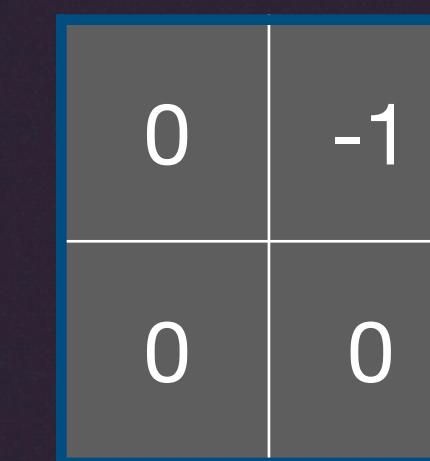
Sobel
Vertical
filter



Horizontal
filtered
image



Vertical
filtered
image



Pooling Layer

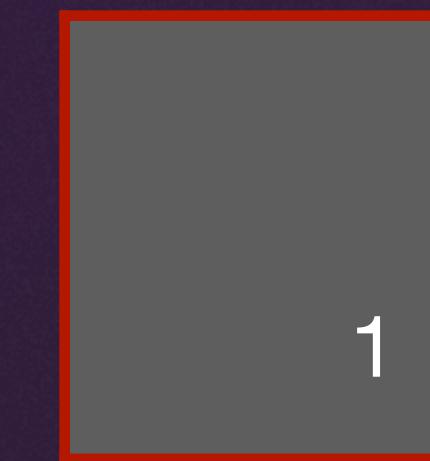
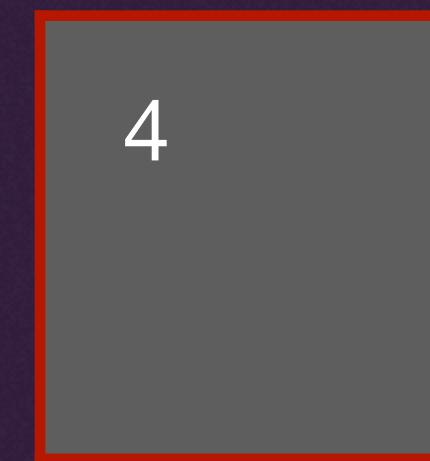
Max Pooling

-1	-2	-1
0	0	0
1	2	1

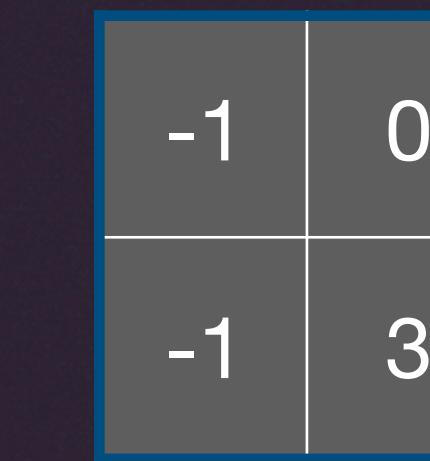
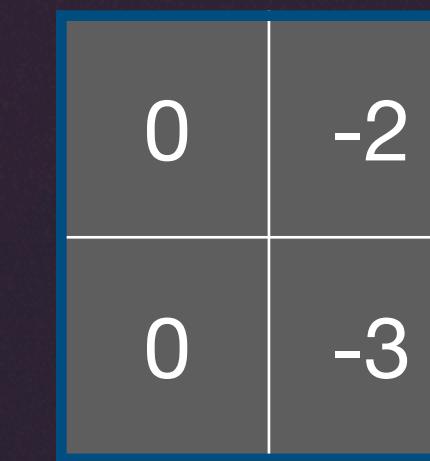
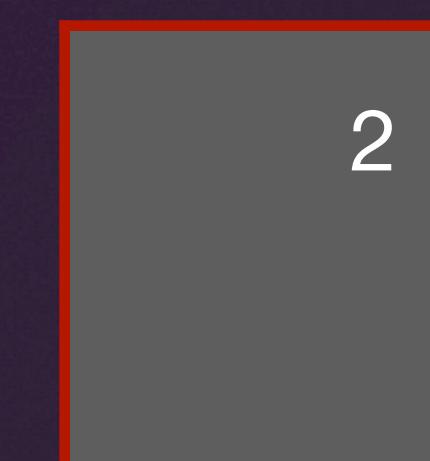
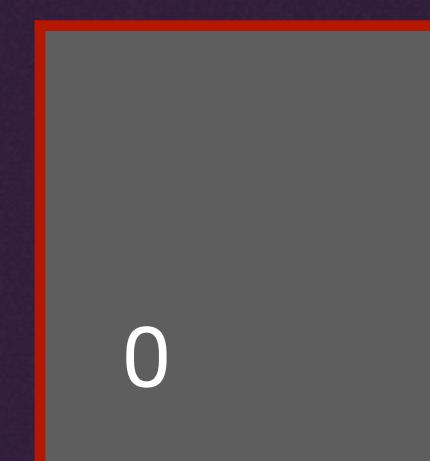
Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

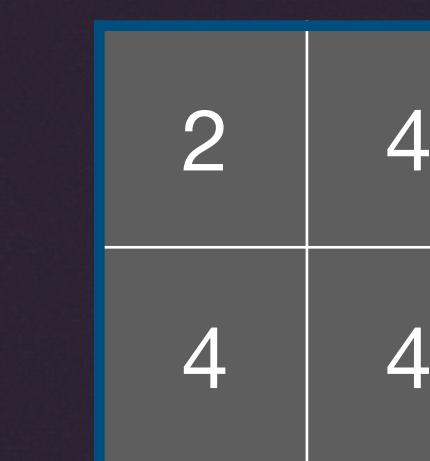
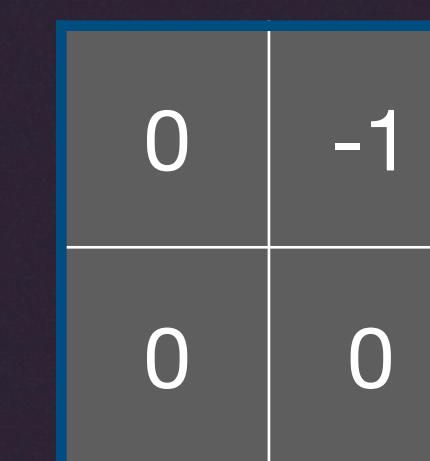
Sobel
Vertical
filter



Horizontal
filtered
image



Vertical
filtered
image



Pooling Layer

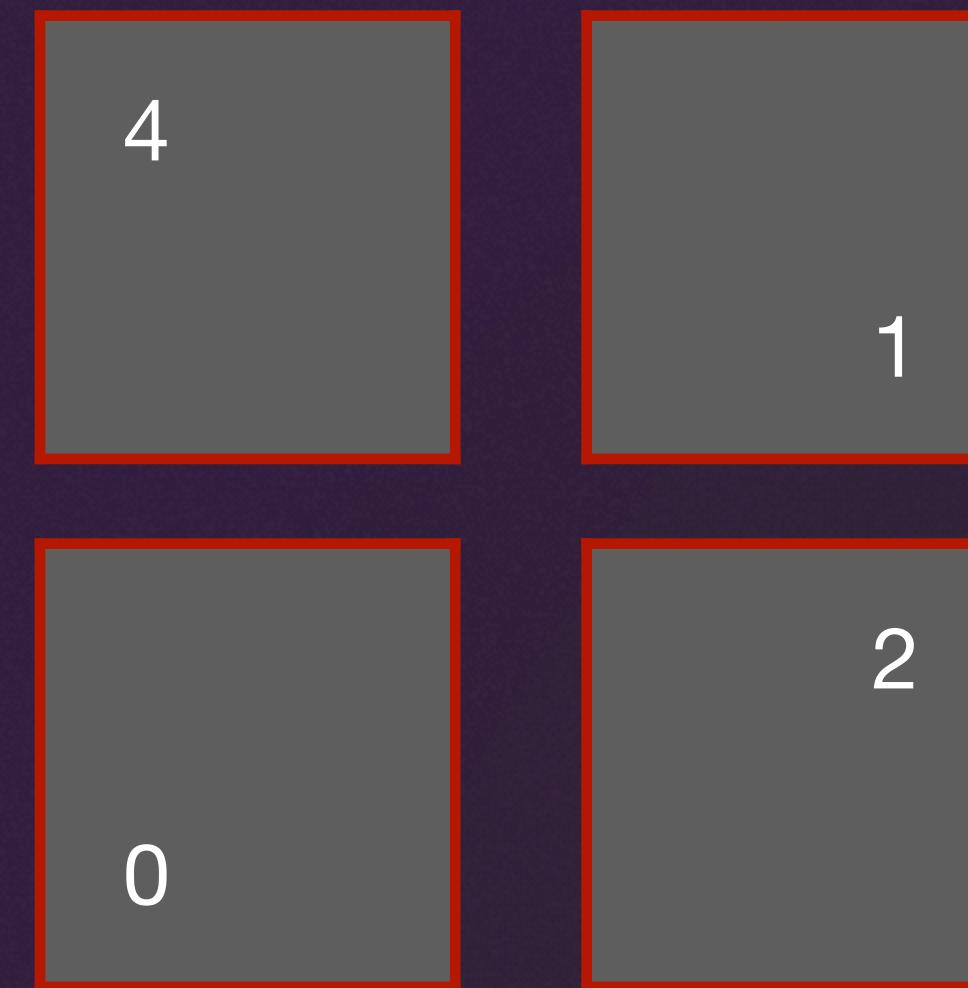
Max Pooling

-1	-2	-1
0	0	0
1	2	1

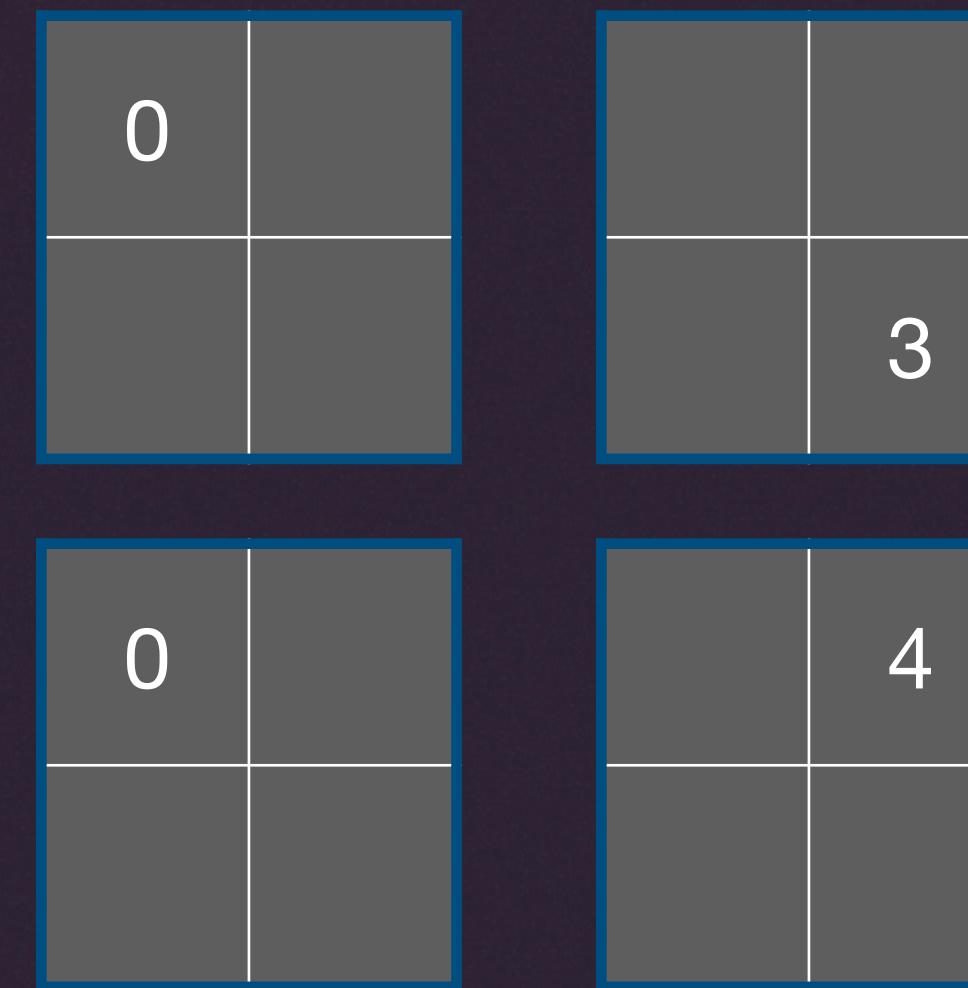
Sobel
Horizontal
filter

-1	0	1
-2	0	2
-1	0	1

Sobel
Vertical
filter



Horizontal
filtered
image

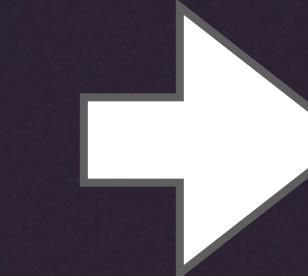
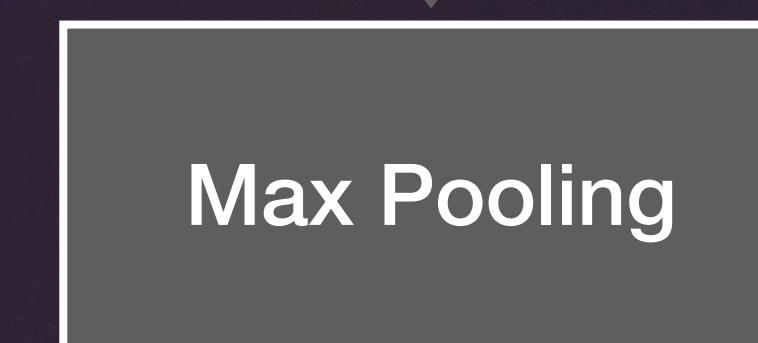
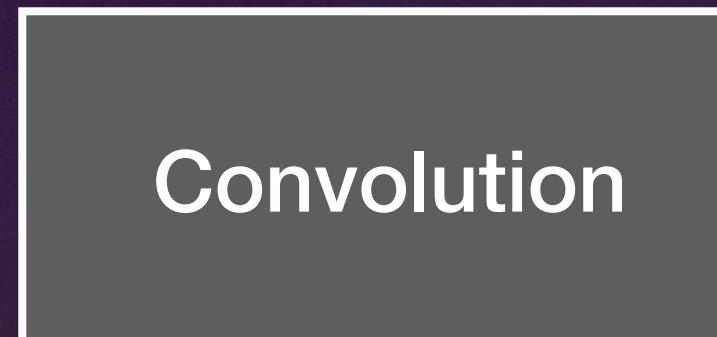
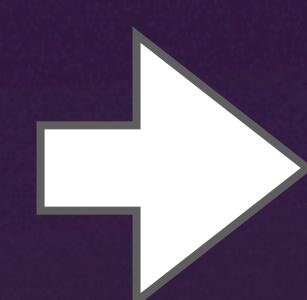


Vertical
filtered
image

Pooling Layer

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

6 × 6 image



New image, but
smaller

4	1
0	3
0	4

2 × 2 image
Each filter is a channel

2 Dimensional Convolution

QUIZ TIME

2 Dimensional Convolution

QUIZ TIME

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

1	0	0
0	1	0
0	0	1

A. 0

B. 1

C. 2

D. 3

2 Dimensional Convolution

QUIZ TIME

0	0	0	1	0	0
1	1	1	0	1	0
1	1	1	0	0	1
0	0	0	0	1	1
0	0	0	0	1	1
0	0	0	0	1	1

-1	1	-1
-1	1	-1
-1	1	-1

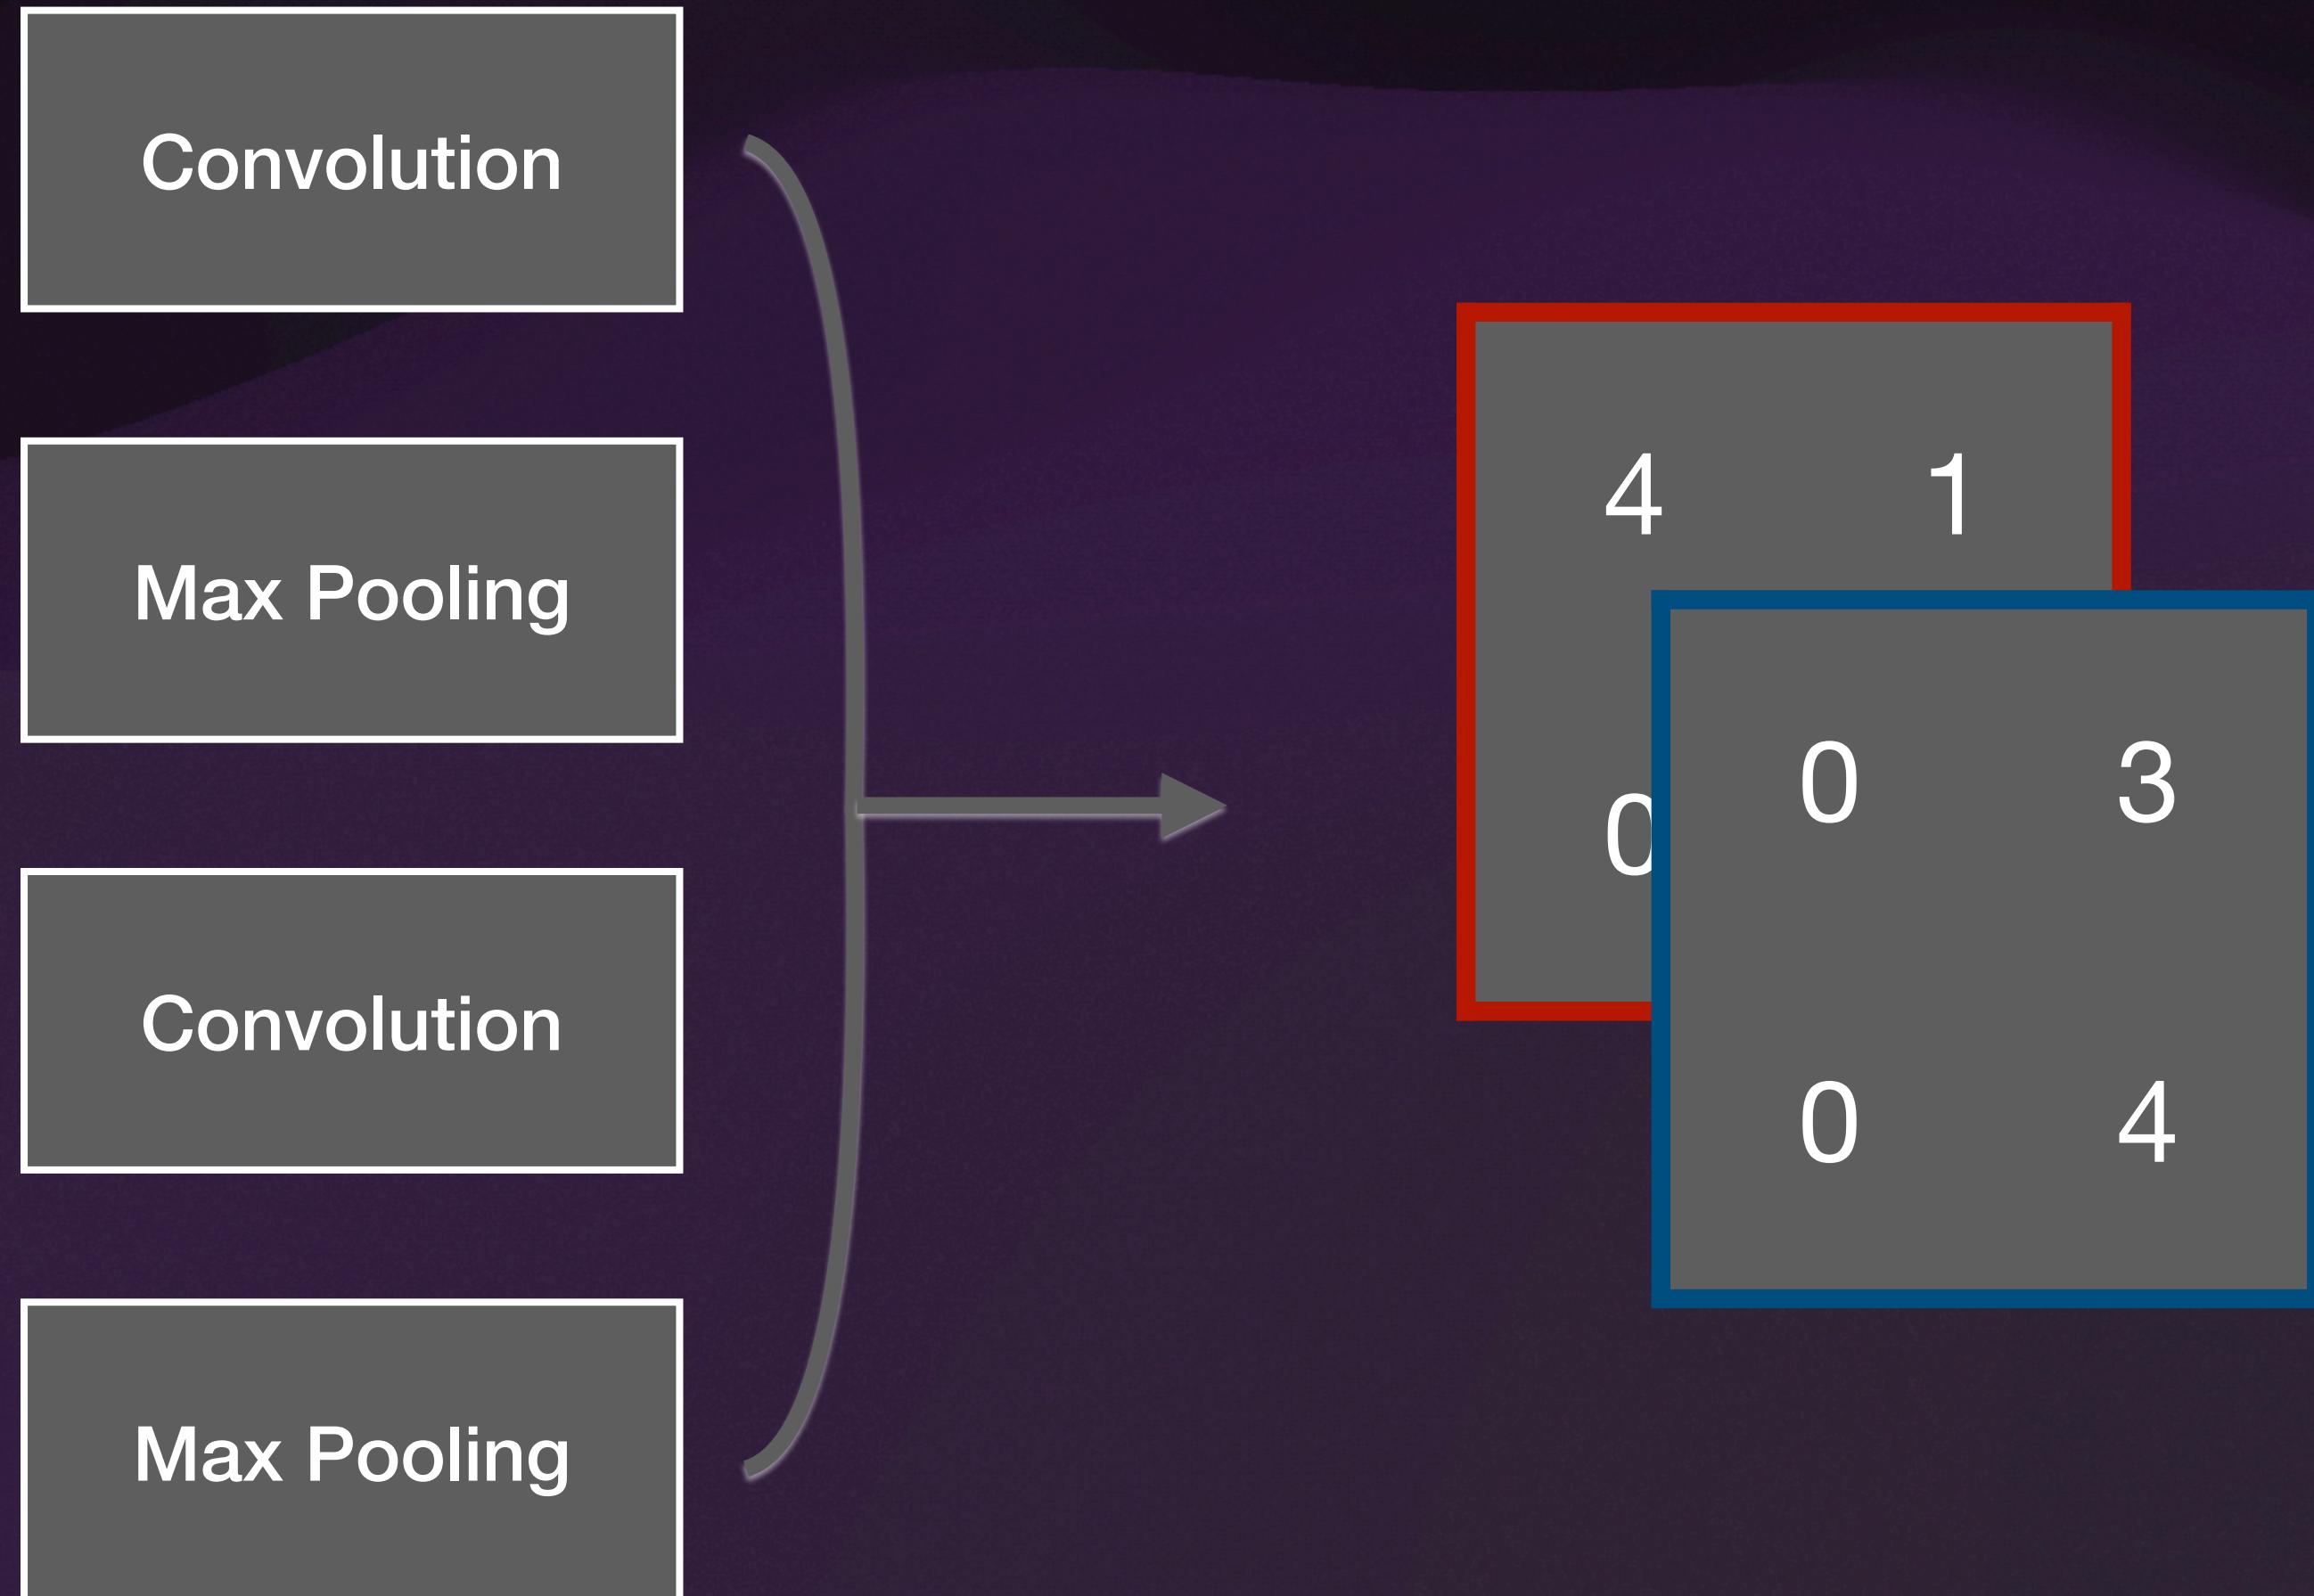
A. 0

B. 1

C. 2

D. 3

Multiple Convolution Layers

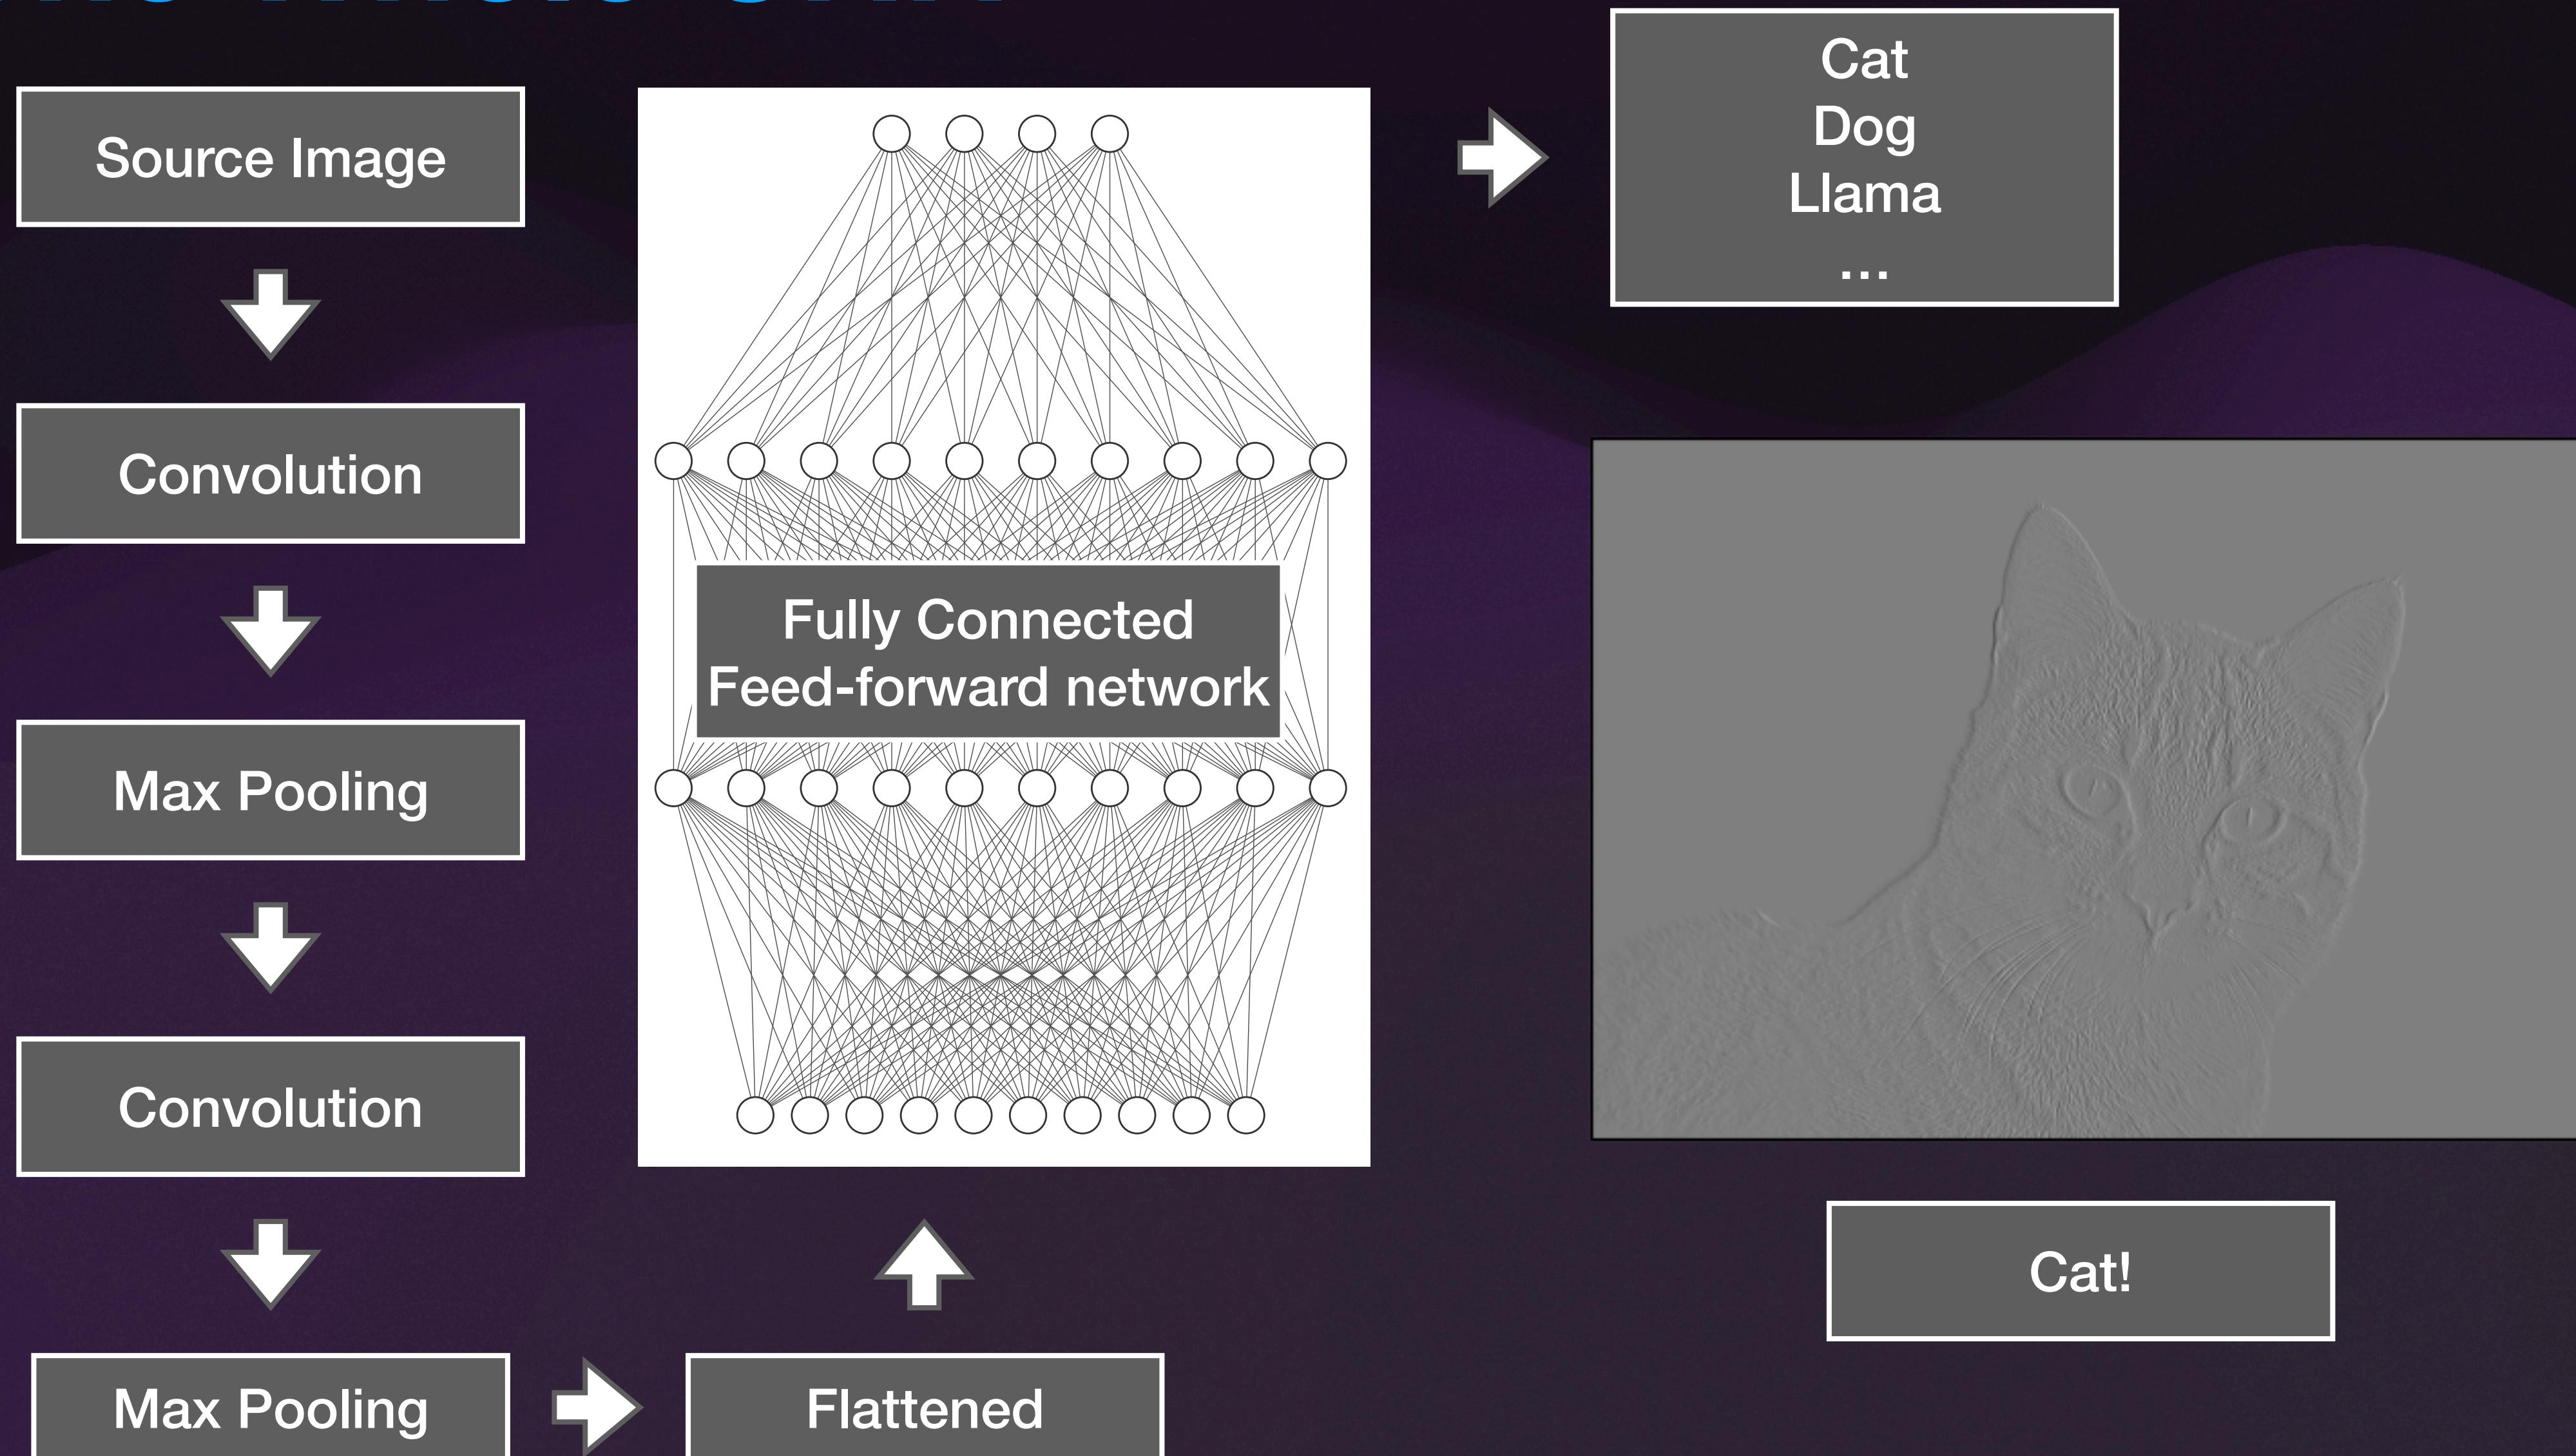


Smaller than the original image

The number of channels is the number of filters

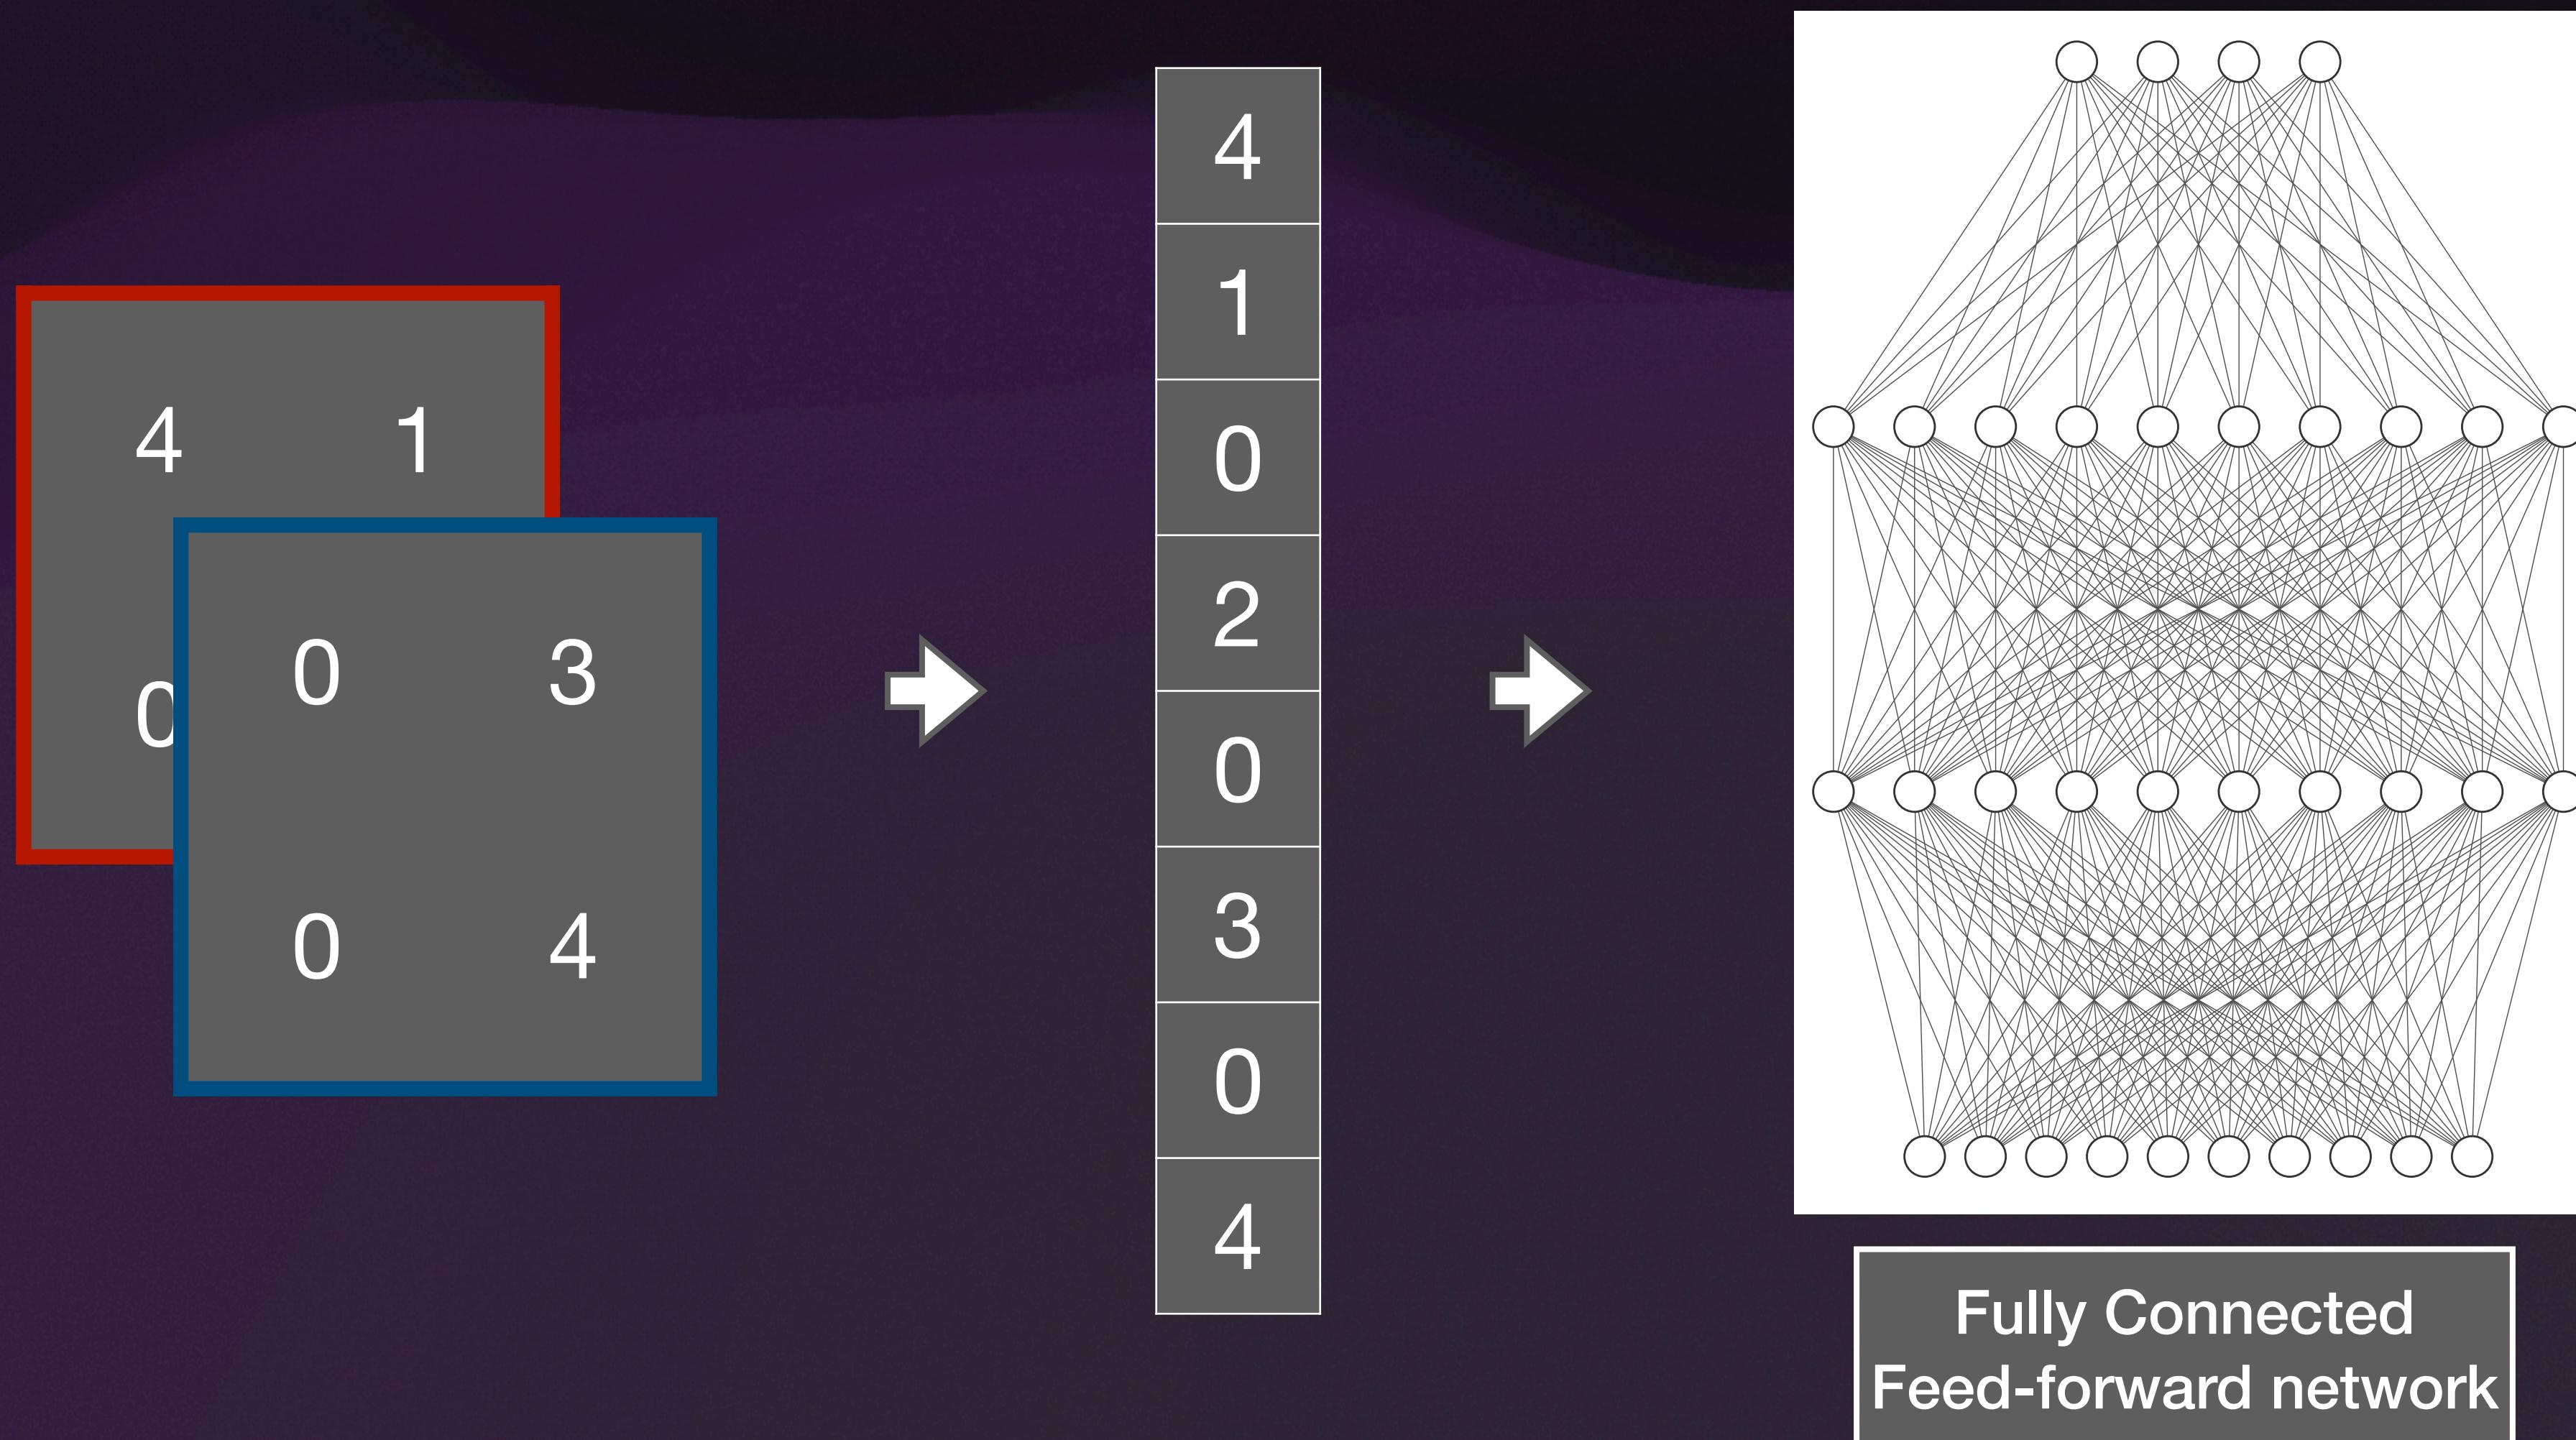
Subsequent layers look for combinations of features from previous layers!

The Whole CNN



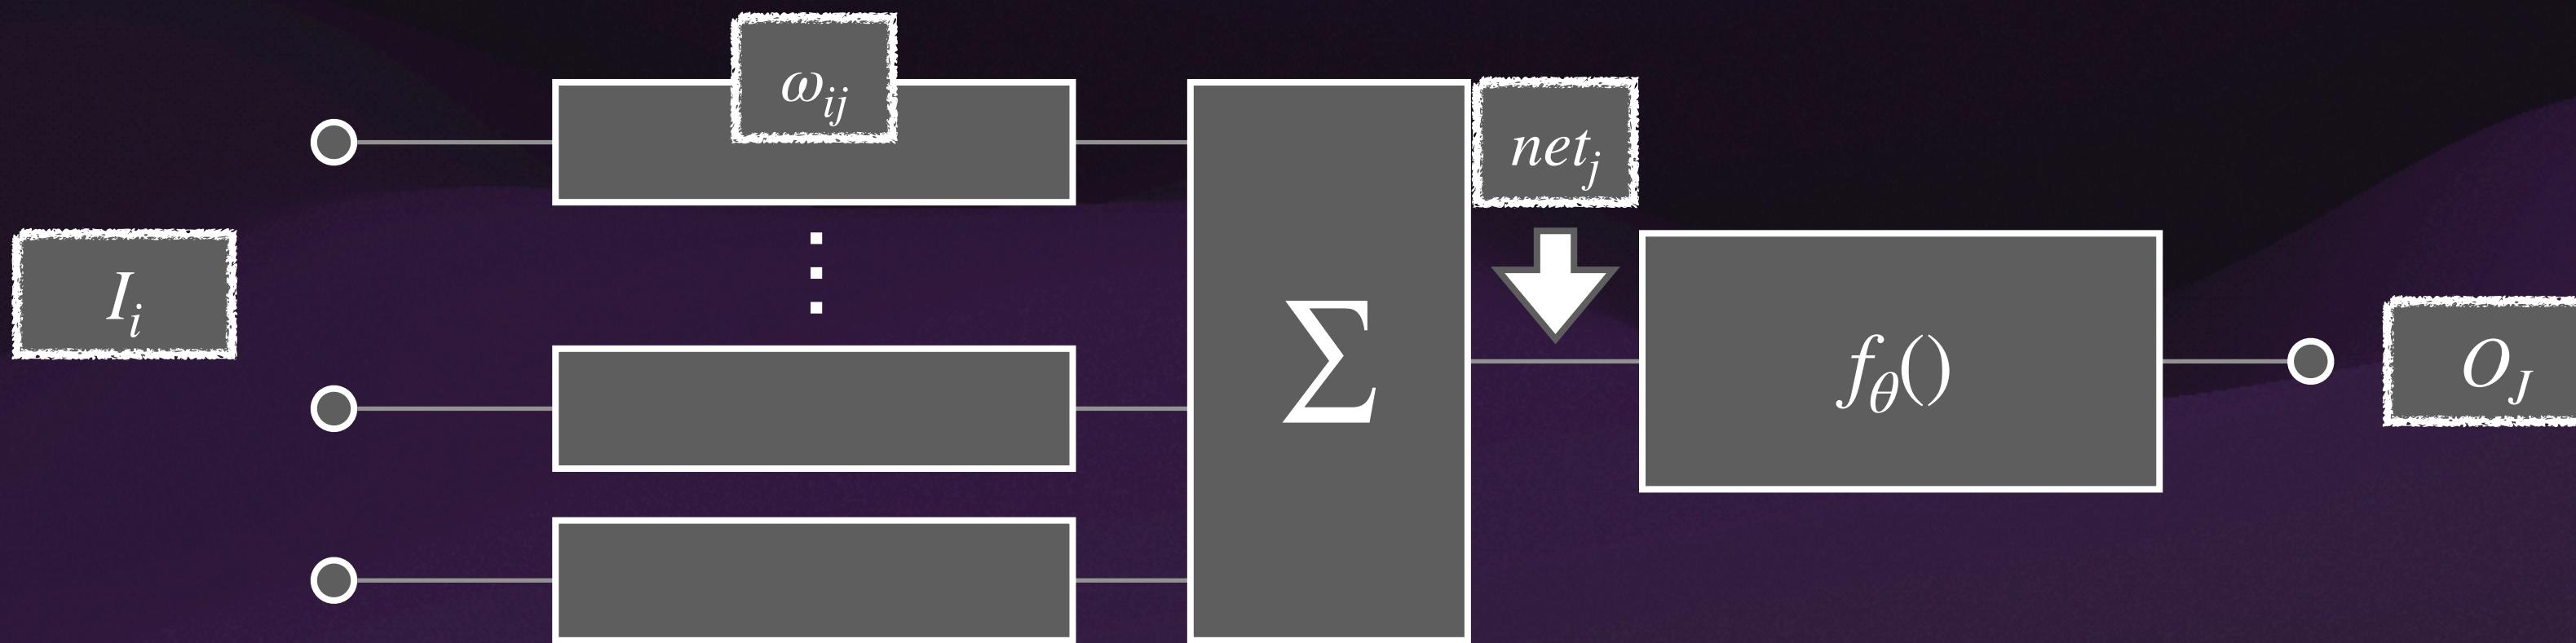
Cat image by Alvesgaspar, Wikimedia Commons

Flattening



Neurons

The McCulloch and Pitts Neuron (1943)



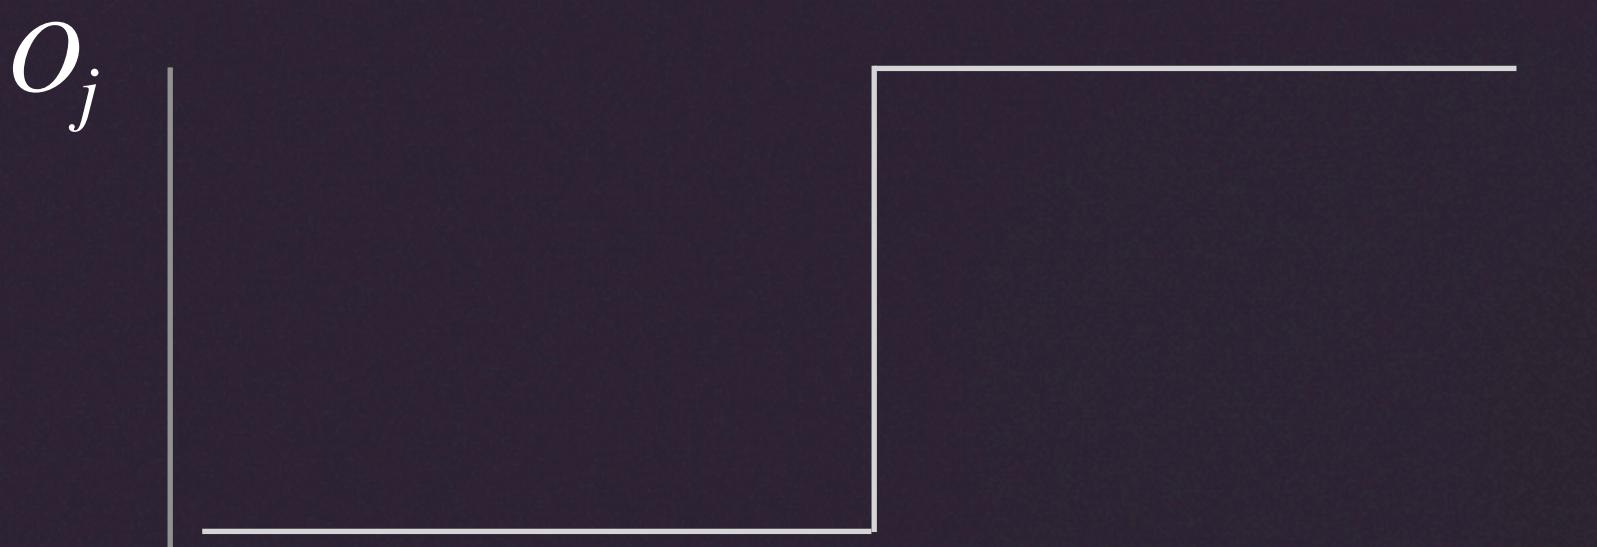
$$O_j = f_\theta(net_j)$$

$$net_j = \sum_i \omega_{ij} i_i$$

Where ω_{ij} is the weight, θ is the threshold, and:

$$f_\theta(x) = 1 \text{ if } x \geq \theta$$

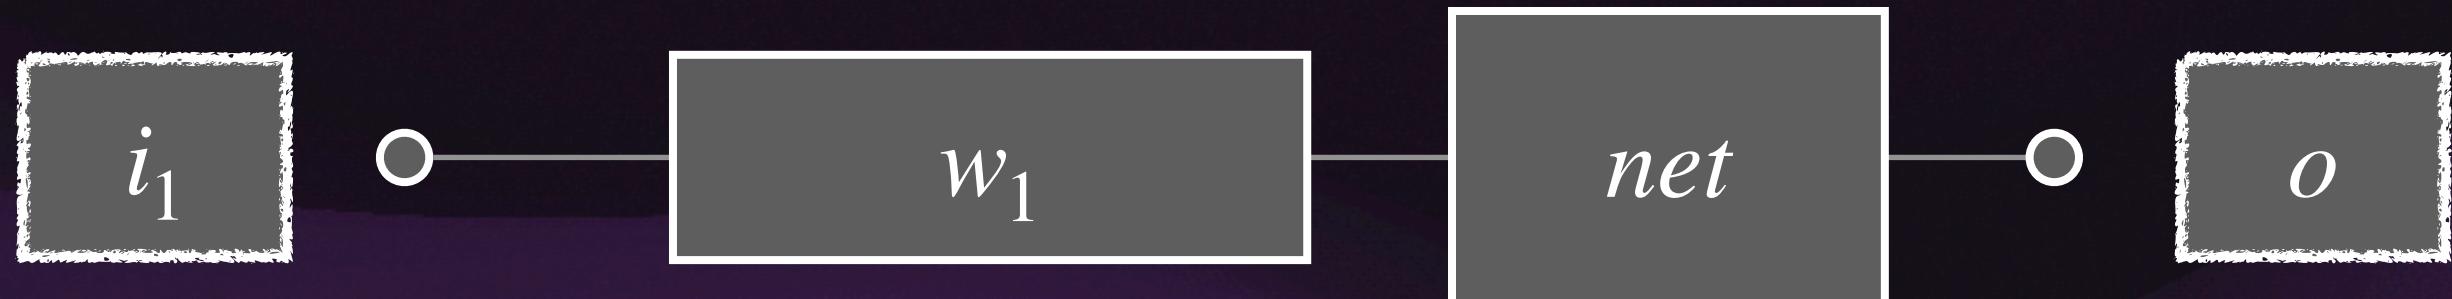
$$f_\theta(x) = 0 \text{ if } x < \theta$$



Example step activation
function $f_\theta(x)$

Single Input Neuron

NB: i_1 not Boolean in this example



$$net = i_1 \times w_1$$
$$o = \begin{cases} 1, & \text{if } net \geq \theta \\ 0, & \text{if } net < \theta \end{cases}$$

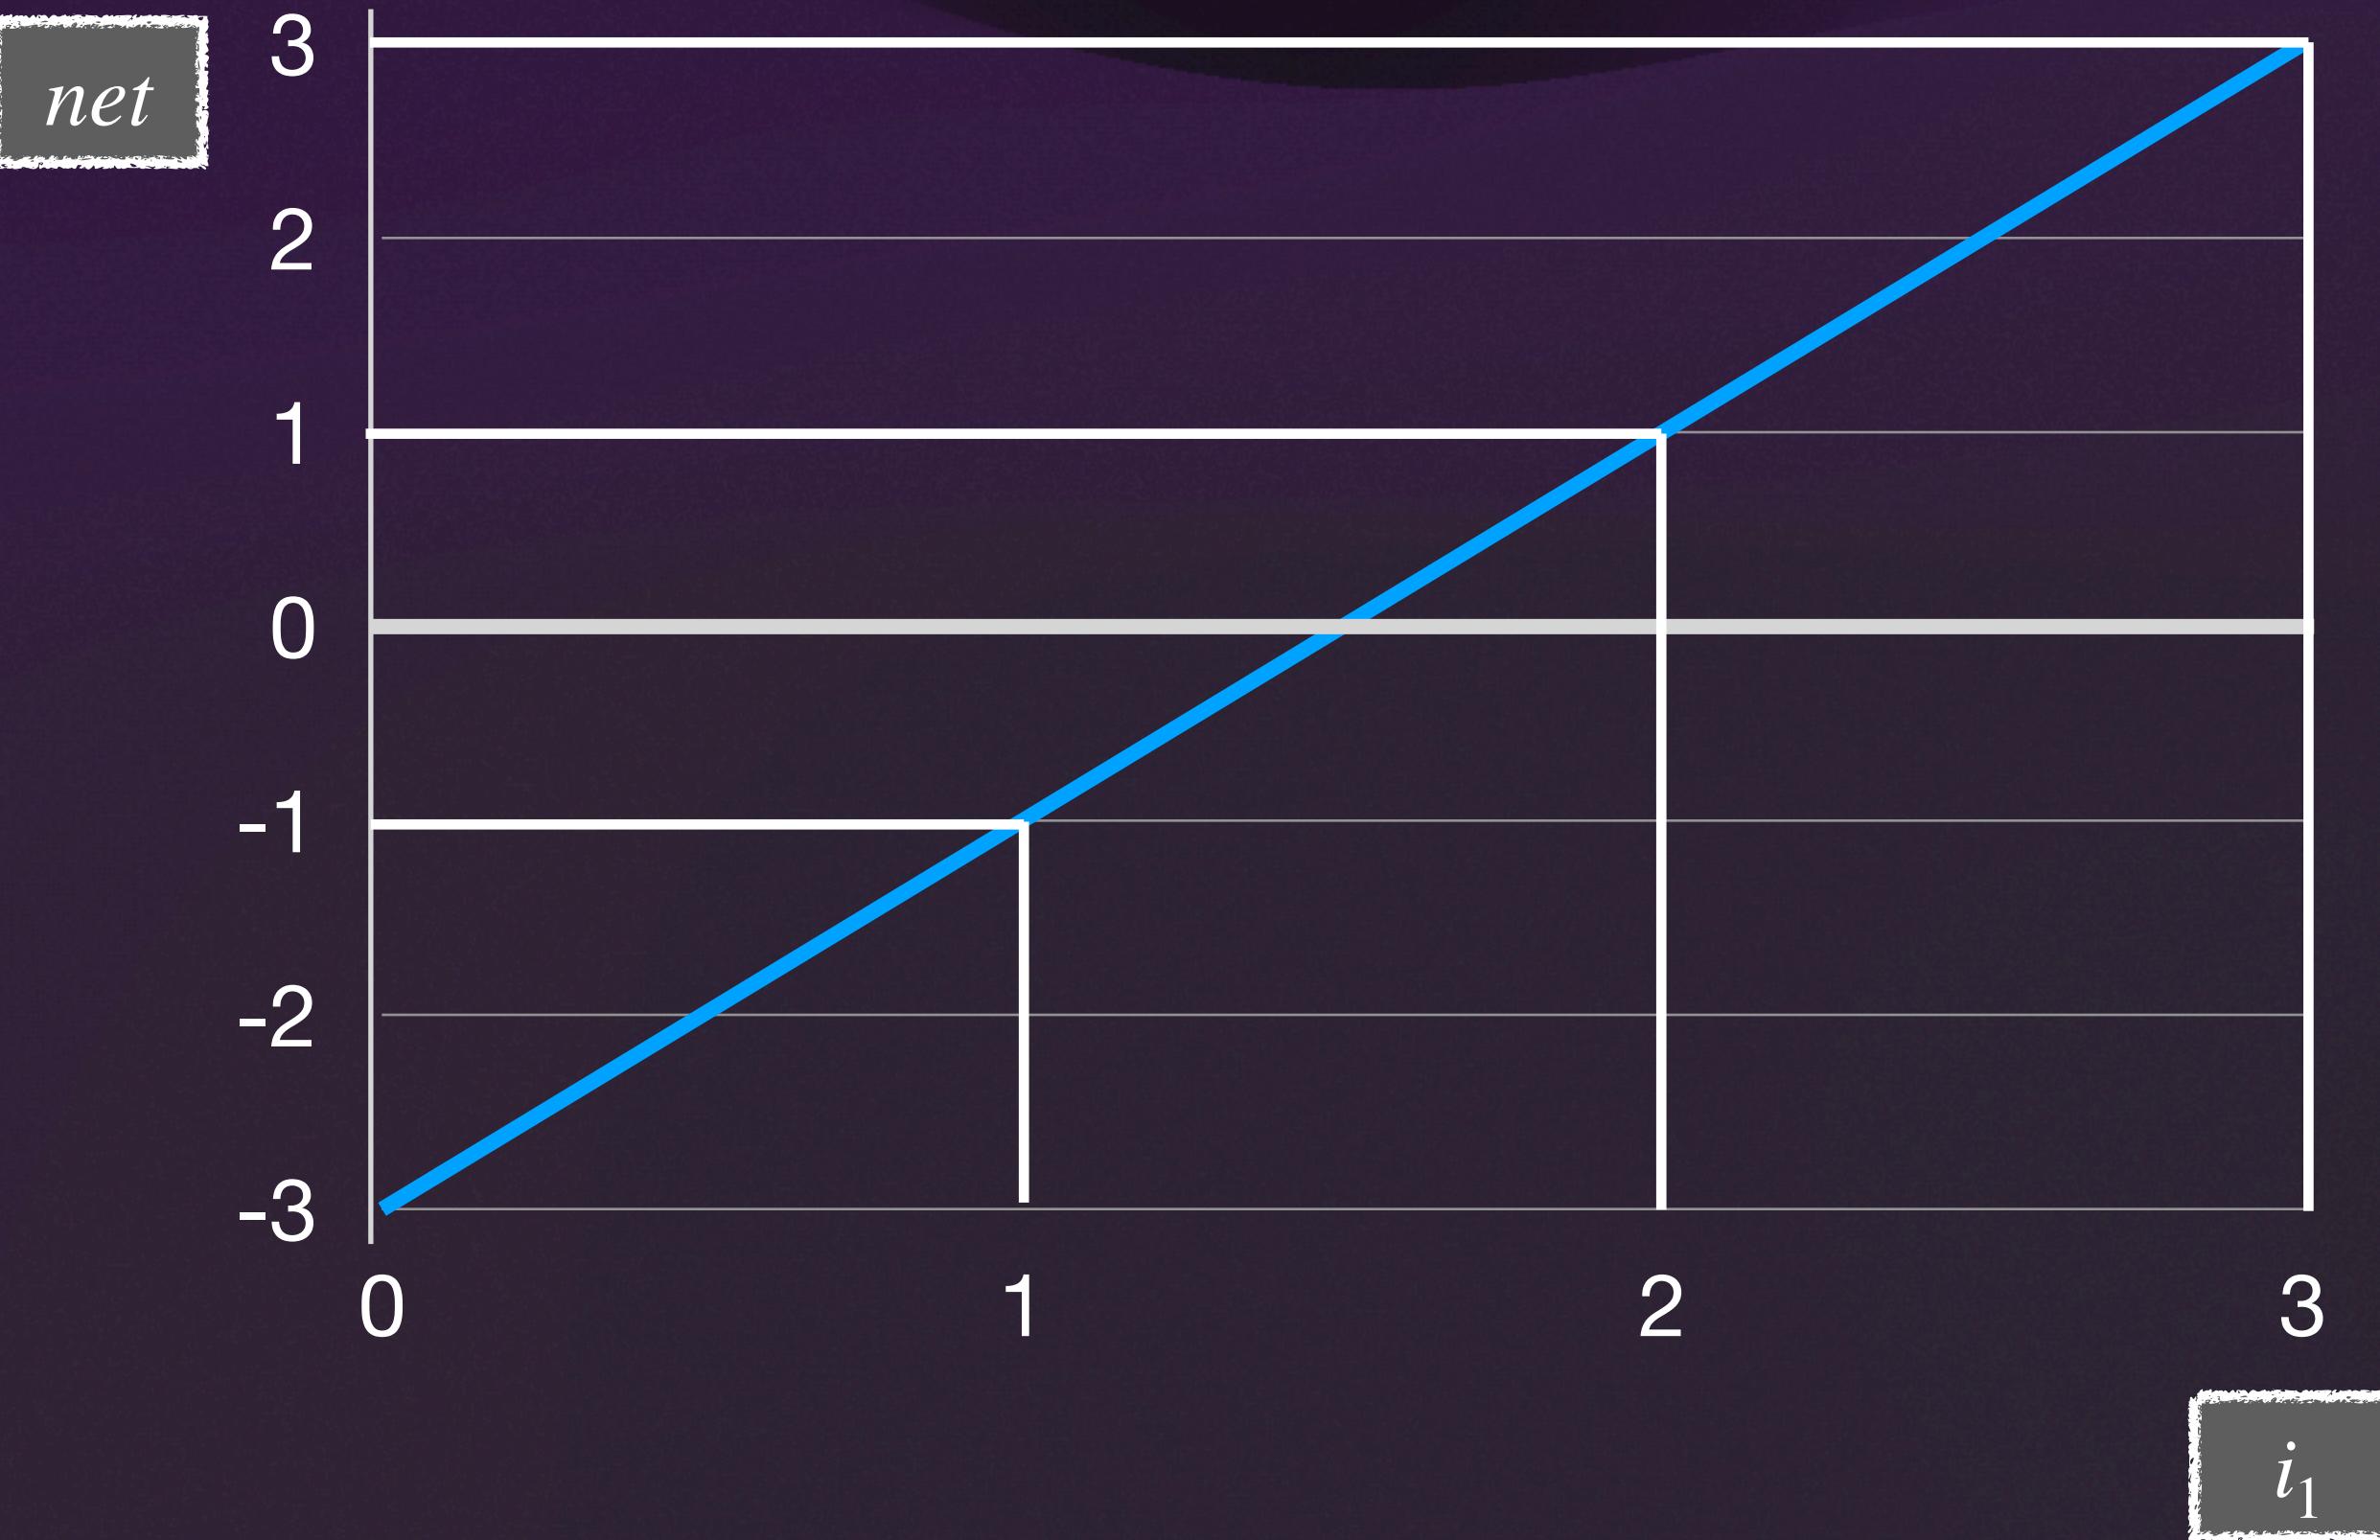
For threshold θ , may rewrite as:

$$net = i_1 \times w_1 - \theta$$
$$o = \begin{cases} 1, & \text{if } net \geq 0 \\ 0, & \text{if } net < 0 \end{cases}$$

Analysis

- Assume $w_1 = 2$ and $\theta = 3$

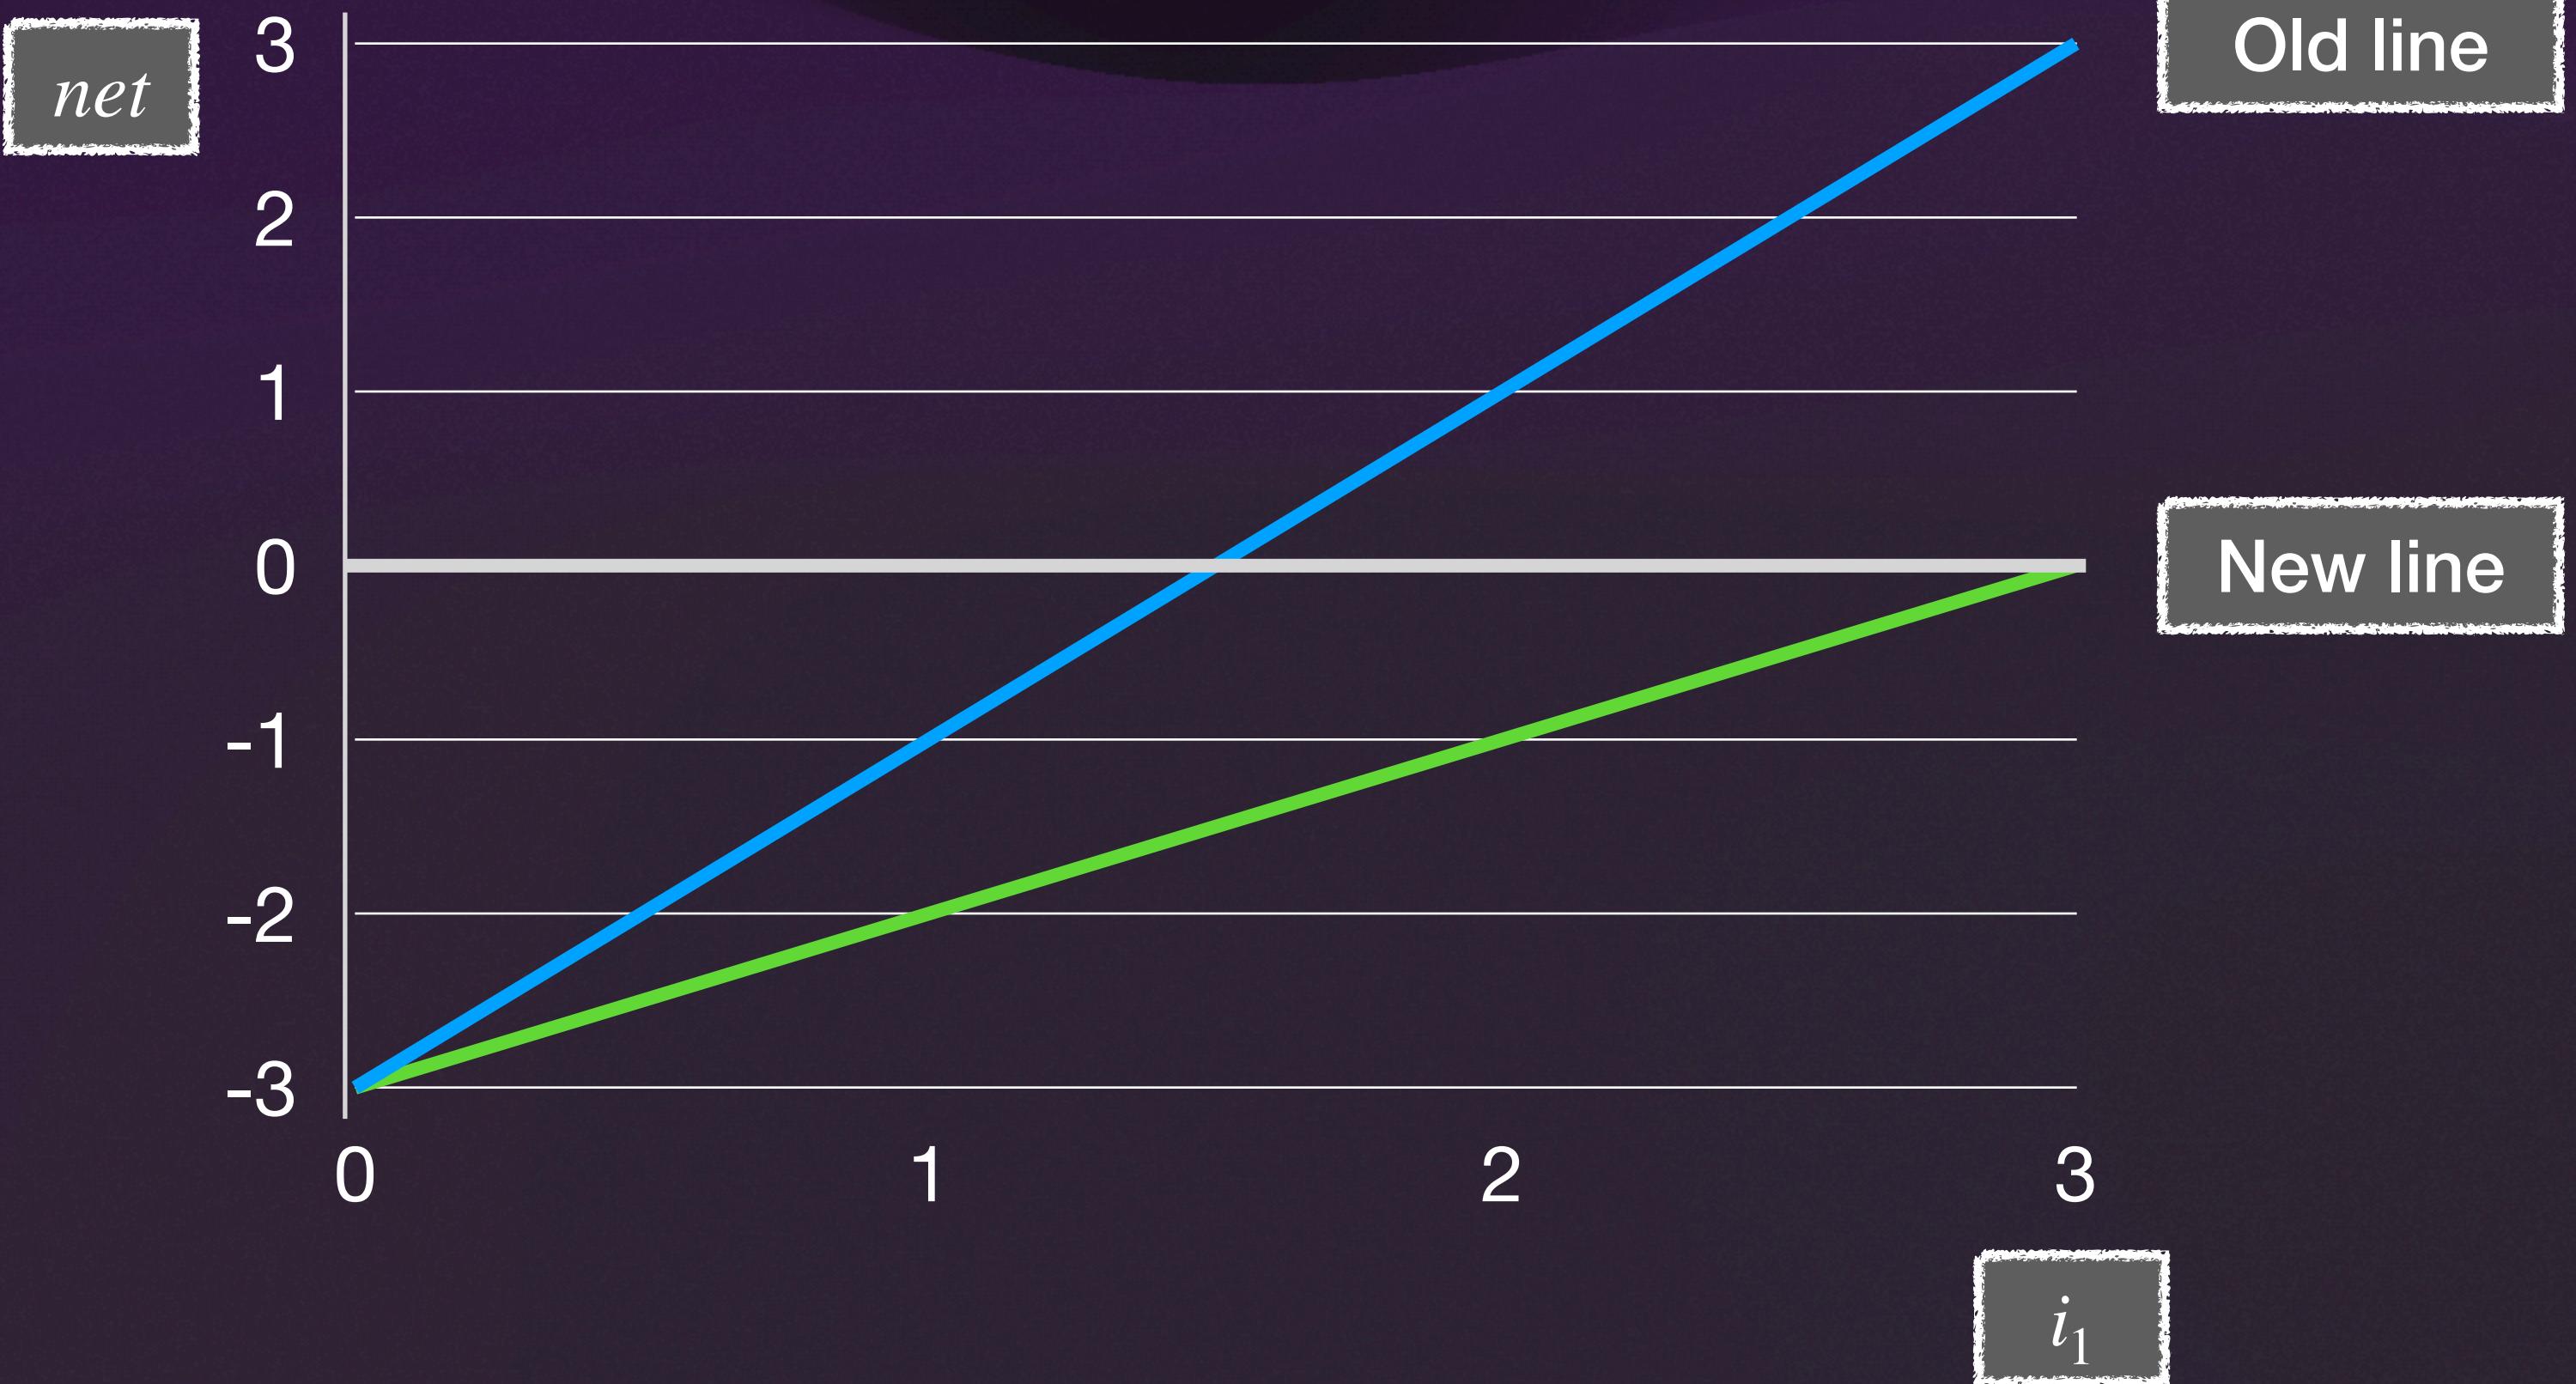
i_1	net
0	-3
1	-1
2	1
3	3



Analysis

- Changing w_1 changes gradient e.g. $w_1 = 1$

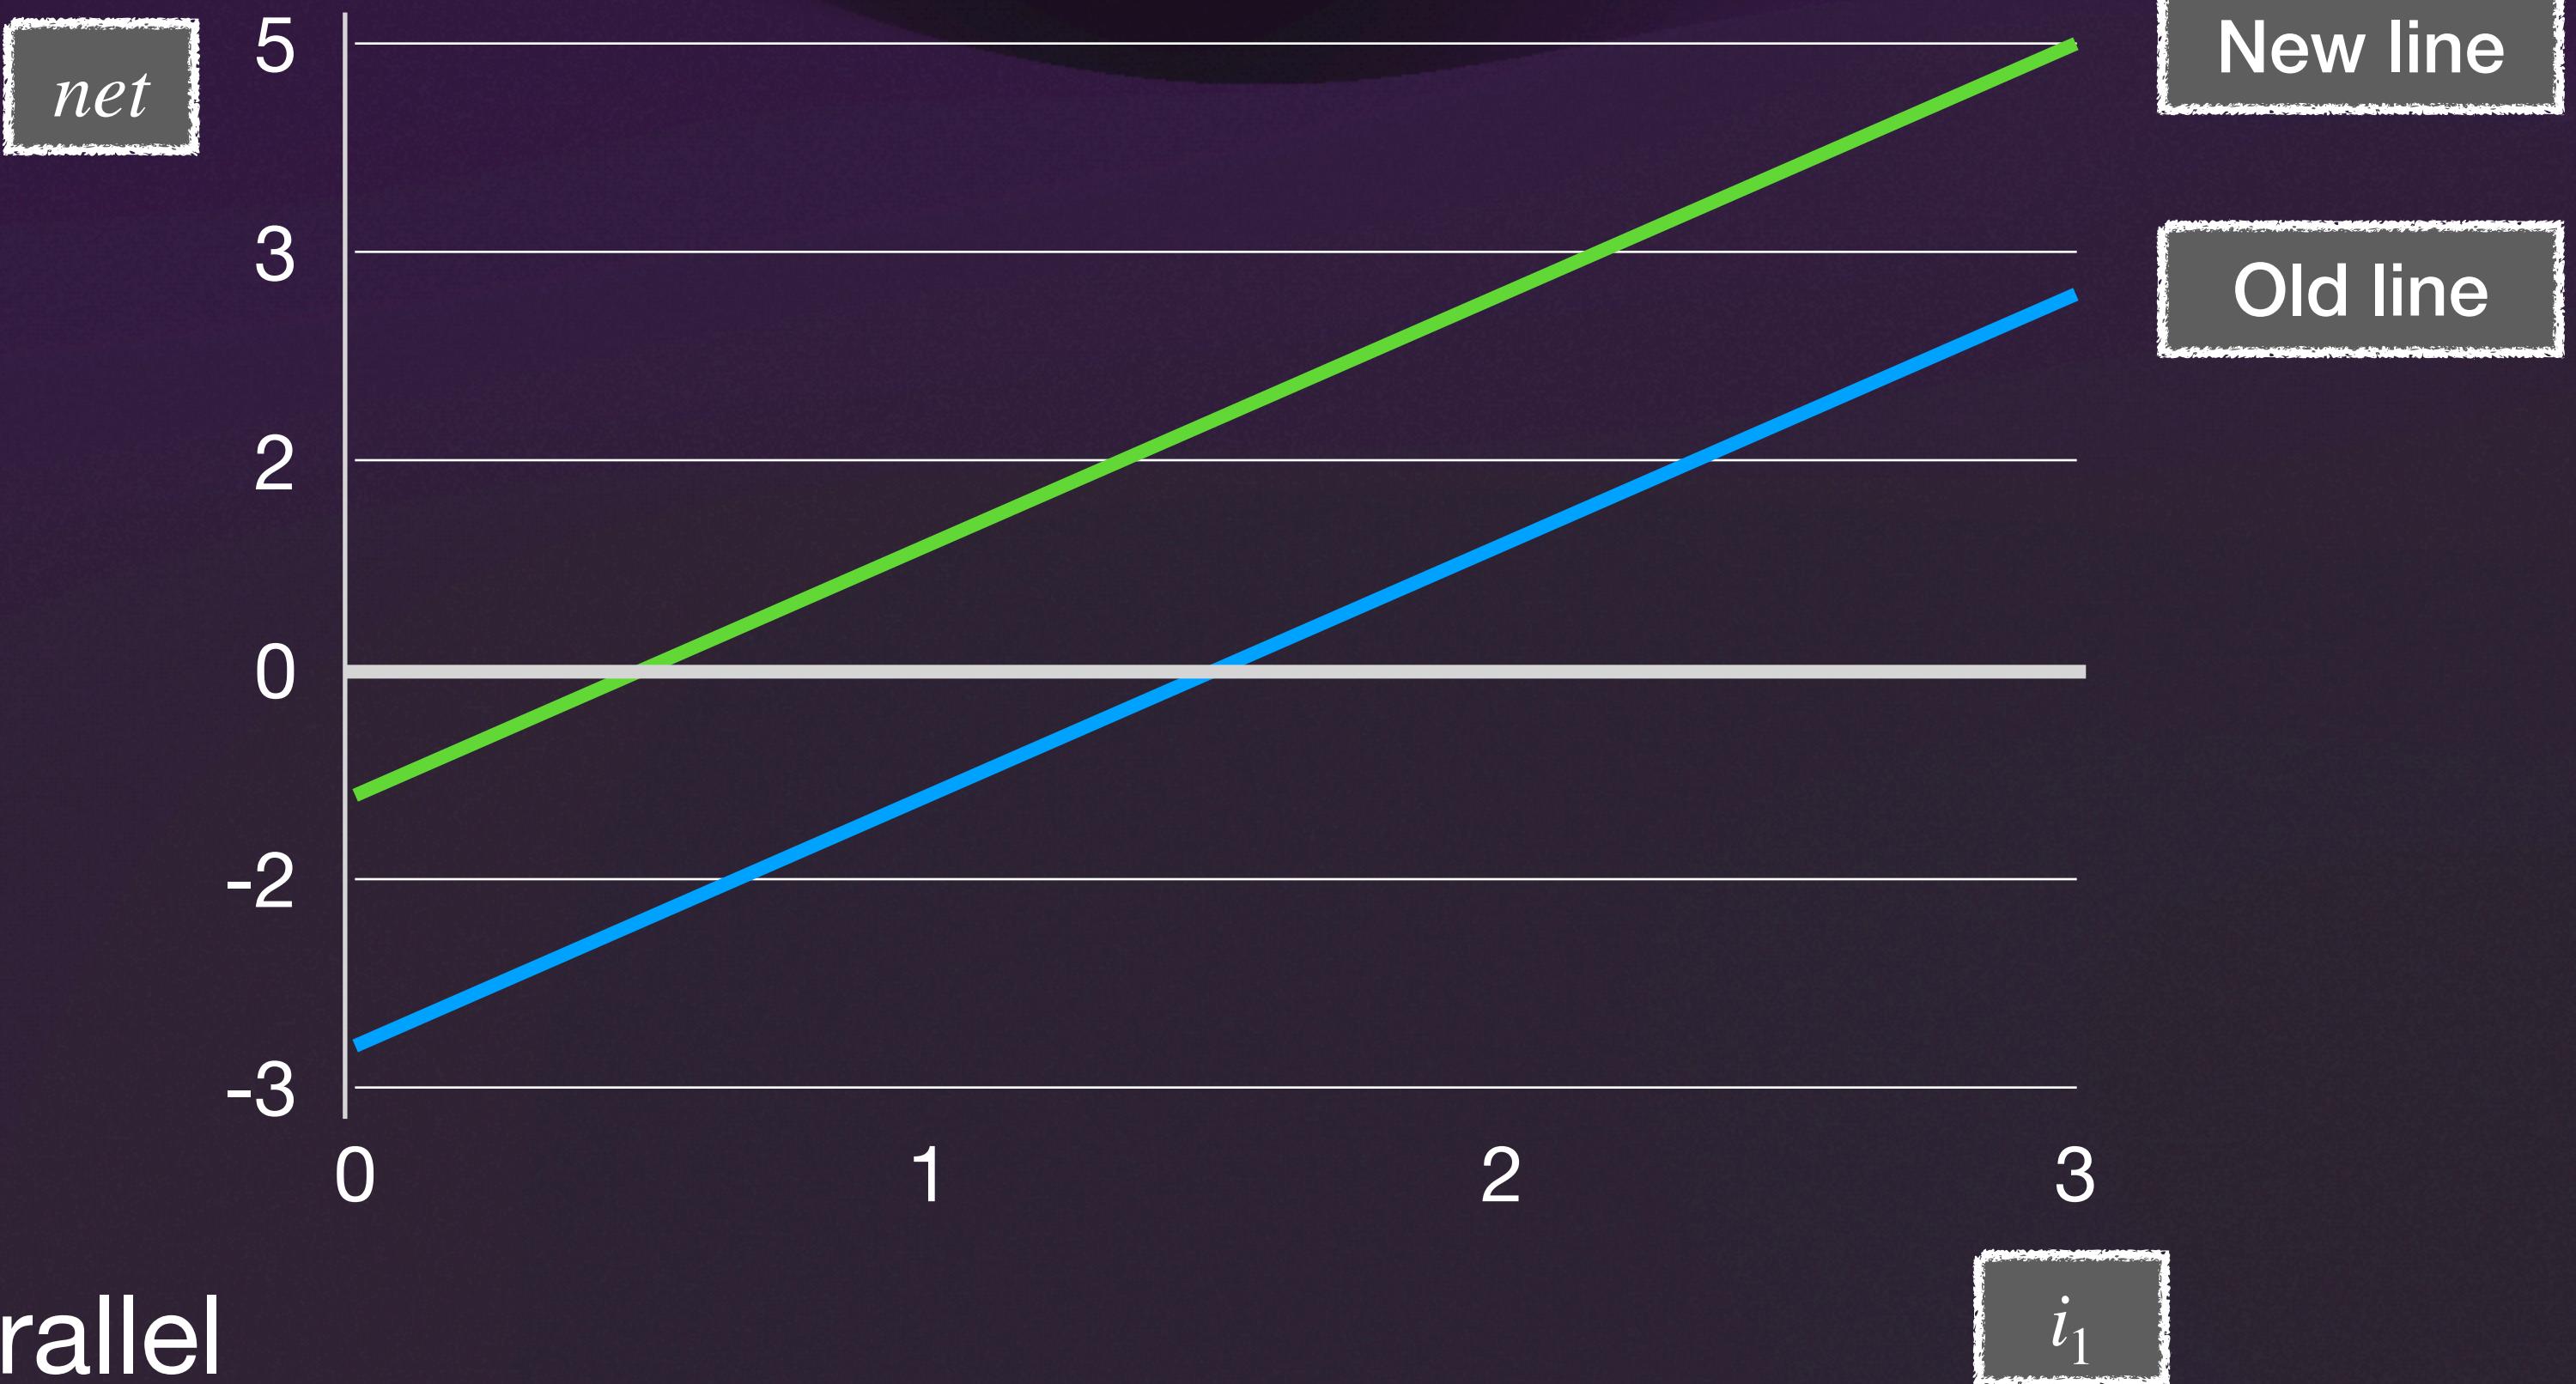
i_1	net
0	-3
1	-2
2	-1
3	0



Analysis

- Changing θ changes the *net* axis intercept e.g. $\theta = 1$

i_1	<i>net</i>
0	-1
1	1
2	3
3	5

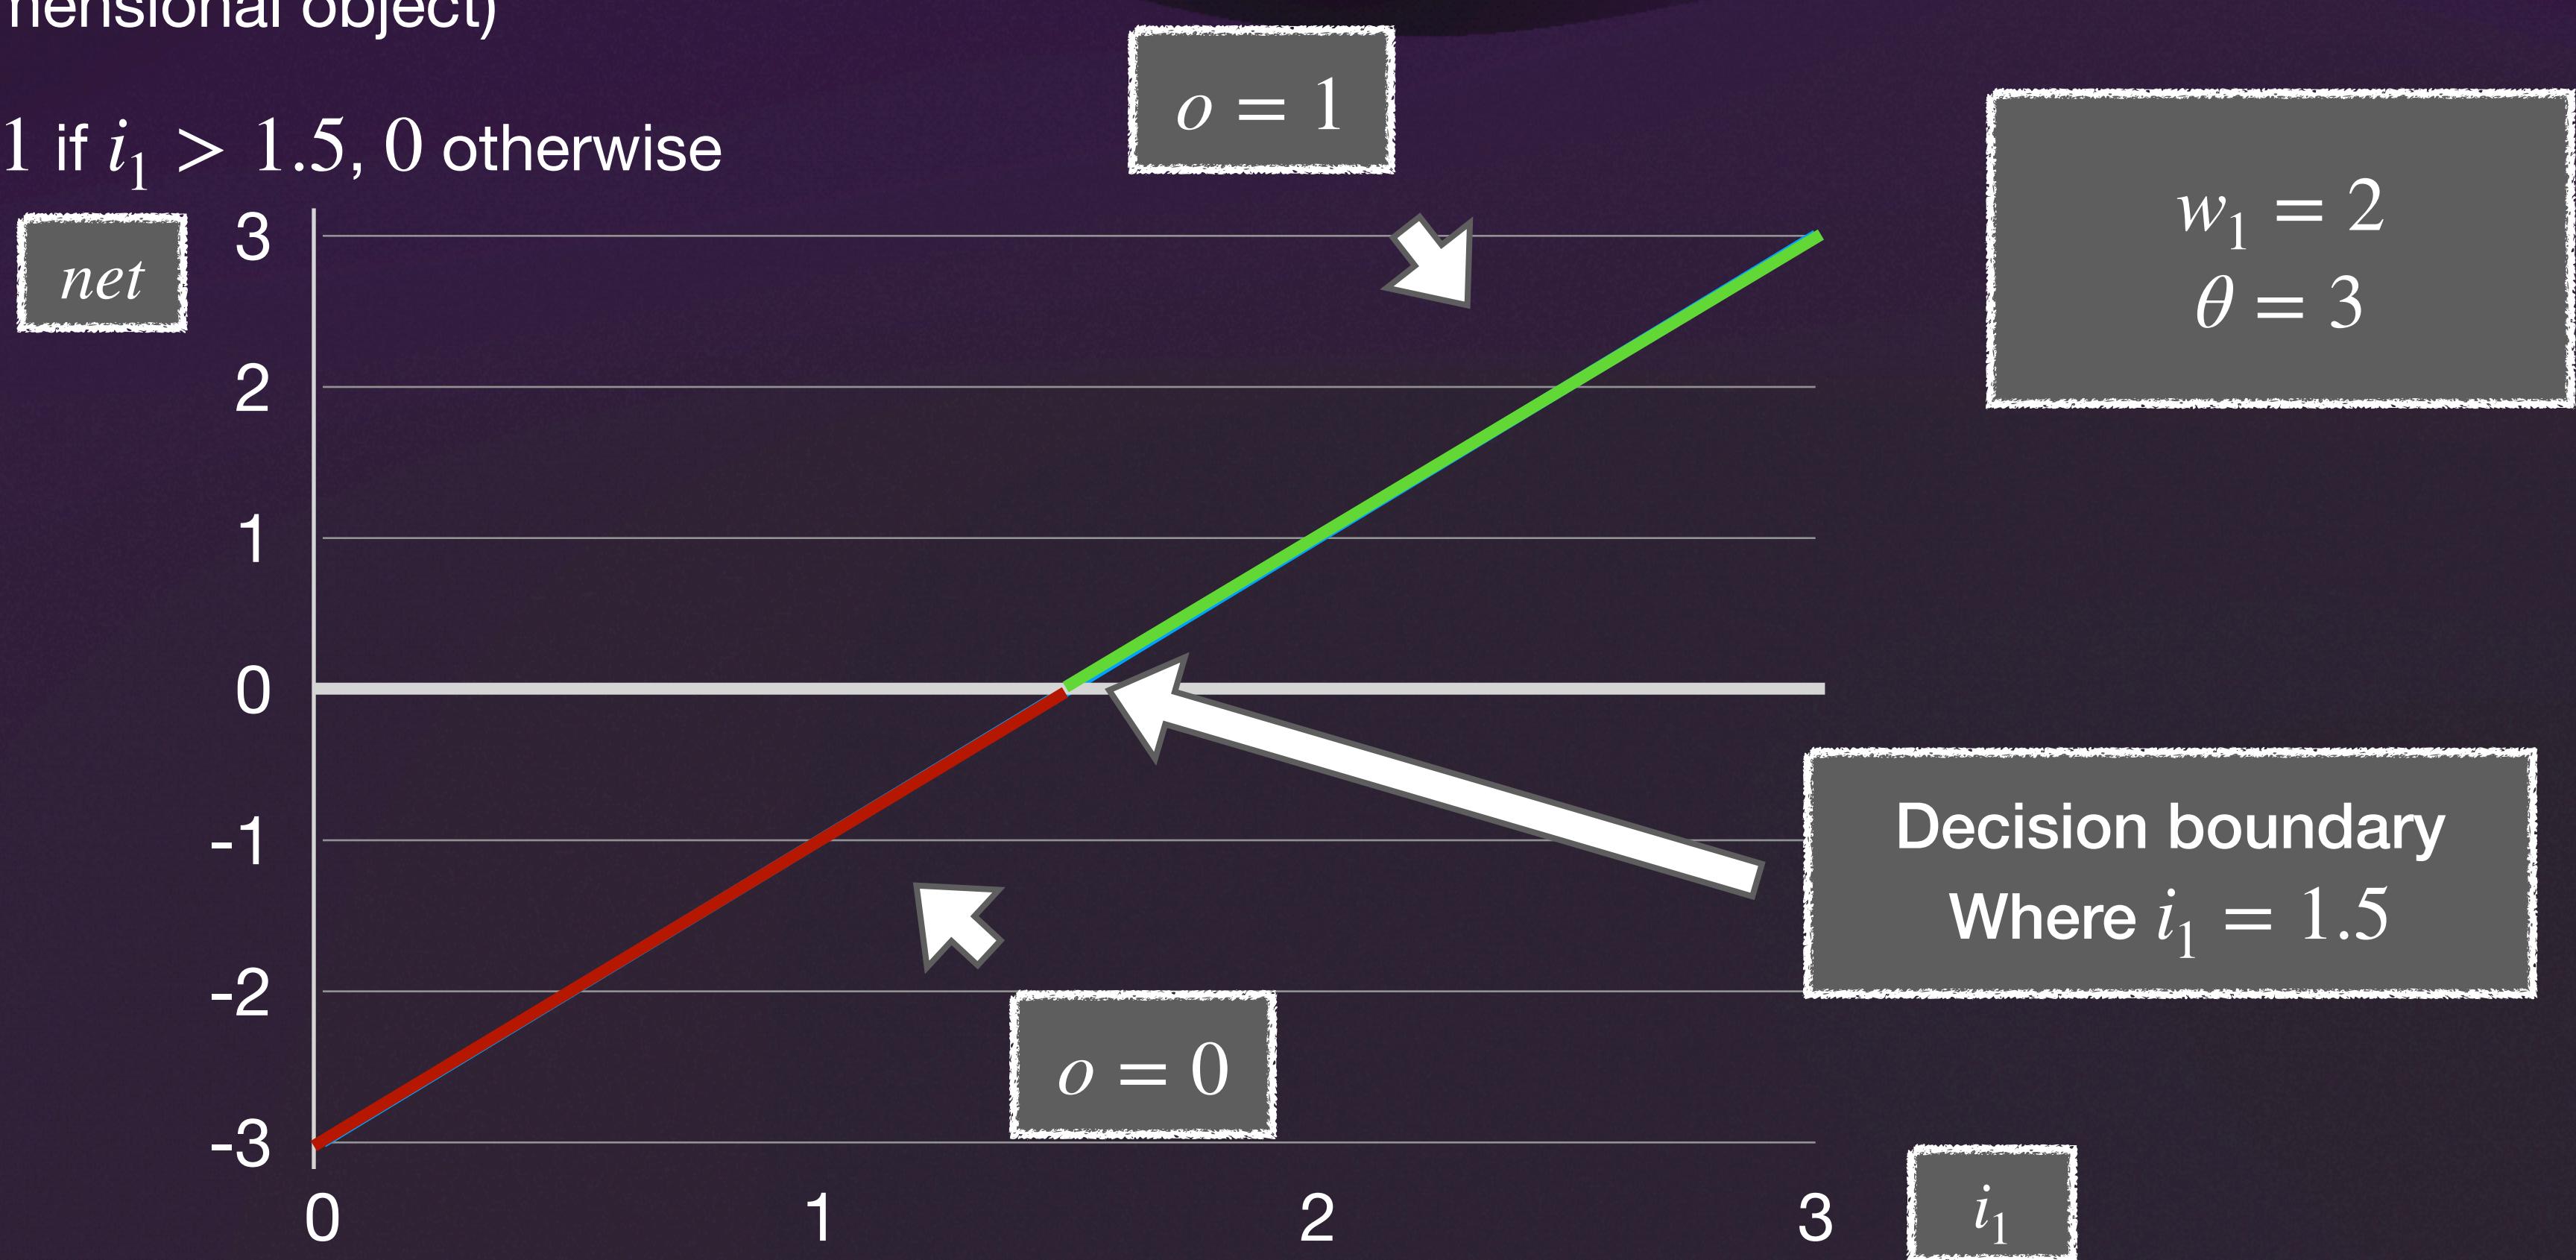


- Old and New lines are parallel

Analysis

- Decision Boundary: Where straight line formed by *net* equations meets straight line formed by i_1 axis
 - It is a single point (0 dimensional object)
 - Original example: $o = 1$ if $i_1 > 1.5$, 0 otherwise

i_1	<i>net</i>
0	-3
1	-1
2	1
3	3



Limitations

- Consider again a neuron with one input. It takes input x , multiplies it by a weight m , and tests if it exceeds threshold c .

$$z = mx - c$$

- It outputs 1 if $z > 0$, otherwise 0
- It is a one-dimensional object that divides a 2 dimensional space
- Although a neuron can divide a space, it can *only divide any space by a straight line*

Summary

- Convolution: Allows the extraction of spatial features in data
- Pooling: Shrink the data
- Flattening: Make the data convenient for a network
- Individual Neurons: Allow splitting a decision space in 2

