

CNN example-Skeleton

March 3, 2025

This example is similar to, although not identical to, examples given in Chapter 7 of the book [Deep Learning with Python, Second Edition](#).

1 Convolutional Neural Network example

This second example looks at generating a convolutional neural network. Again, its function is both to investigate the network topology and its abilities but also to familiarise you with how to construct such a network in KERAS.

This example will again use the functional API components of KERAS

The network will again use a character recognition task using the standard MNIST dataset.

2 Defining the network architecture

As for the first network, this is the first section you need to write yourselves.

The workshop script takes you through what you need to do. Note again the Inputs are defined for you however. Your function is to fill in the missing sections to define the layers you will need to employ. Naturally in this section you will need to define convolutional, pooling and flattening layer (the latter is already included).

```
[1]: from tensorflow import keras
    from tensorflow.keras import layers

    inputs = keras.Input(shape=(28, 28, 1))
    # First convolutional layer
    x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(inputs)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)

    # Second convolutional layer
    x = layers.Conv2D(64, (3, 3), activation="relu", padding="same")(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)

    # Third convolutional layer (newly added)
    x = layers.Conv2D(128, (3, 3), activation="relu", padding="same")(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)

    # Flatten the output before feeding into Dense layers
```

```

x = layers.Flatten()(x)
x = layers.Dense(128, activation="relu")(x)

# Output layer (10 classes for digits)
outputs = layers.Dense(10, activation="softmax")(x) # 10 classes for digits

model = keras.Model(inputs=inputs, outputs=outputs)

```

```

2025-02-10 12:43:47.226394: E
tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:9342] Unable to
register cuDNN factory: Attempting to register factory for plugin cuDNN when one
has already been registered
2025-02-10 12:43:47.226507: E
tensorflow/compiler/xla/stream_executor/cuda/cuda_fft.cc:609] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
2025-02-10 12:43:47.226558: E
tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:1518] Unable to
register cuBLAS factory: Attempting to register factory for plugin cuBLAS when
one has already been registered
2025-02-10 12:43:48.802661: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not
find TensorRT

```

Displaying the model's summary

```
[2]: model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 128)	0

flatten (Flatten)	(None, 1152)	0
dense (Dense)	(None, 128)	147584
dense_1 (Dense)	(None, 10)	1290

```
=====
Total params: 241546 (943.54 KB)
Trainable params: 241546 (943.54 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

3 Compiling and training the model

The next stage is to compile the model using an optimiser and an error calculation. Follow the script to deduce what to put here. The images are preprocessed for you as for the first example.

Train the model analogously to the first example.

```
[3]: from tensorflow.keras.datasets import mnist

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype("float32") / 255

model.compile(optimizer="adam",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

model.fit(train_images, train_labels, epochs=5, batch_size=32,
          validation_data=(test_images, test_labels))
```

```
Epoch 1/5
1875/1875 [=====] - 47s 25ms/step - loss: 0.1272 -
accuracy: 0.9602 - val_loss: 0.0465 - val_accuracy: 0.9850
Epoch 2/5
1875/1875 [=====] - 39s 21ms/step - loss: 0.0404 -
accuracy: 0.9875 - val_loss: 0.0432 - val_accuracy: 0.9863
Epoch 3/5
1875/1875 [=====] - 37s 20ms/step - loss: 0.0297 -
accuracy: 0.9910 - val_loss: 0.0264 - val_accuracy: 0.9913
Epoch 4/5
1875/1875 [=====] - 36s 19ms/step - loss: 0.0217 -
accuracy: 0.9931 - val_loss: 0.0266 - val_accuracy: 0.9911
Epoch 5/5
```

```
1875/1875 [=====] - 42s 22ms/step - loss: 0.0175 -  
accuracy: 0.9945 - val_loss: 0.0292 - val_accuracy: 0.9916
```

```
[3]: <keras.src.callbacks.History at 0x7ff70f06faf0>
```

Evaluating the convnet

```
[4]: test_loss, test_acc = model.evaluate(test_images, test_labels)  
print(f"Test accuracy: {test_acc:.3f}")
```

```
313/313 [=====] - 3s 10ms/step - loss: 0.0292 -  
accuracy: 0.9916  
Test accuracy: 0.992
```

Model Type	Architecture	Strengths Weaknesses
MLP (Dense layers only)	3 fully connected layers	Easier to understand, works for simple problems
CNN (Conv + Pool + Dense)	3 Conv layers + Pooling + Dense	Recognizes spatial structures, better for images